The Silicon-Power Gap

A Technical and Financial Analysis of the AI Infrastructure Energy Bottleneck

Lorenzo Frascolla Independent Researcher, Simancas Capital

Version 39.1, 16 October 2026 (version 39 of 6 August 2026 with printing corrections). First circulated May 2026; Section II market snapshot as of June 1, 2026. A public correction log for versions 37 to 39.1 is posted with the paper's supporting files.

Abstract

This paper argues that the binding constraint on AI compute deployment in the next two to four years is not the supply of NVIDIA chips, but the supply of the high-voltage electrical apparatus that steps down, conditions, and distributes utility power for data-center use. Specifically, the large power transformers (primary substation step-down units, plus generator step-up units where on-site generation is involved) and medium-voltage switchgear in the 100 megavolt-ampere class. Using a Black-Scholes-Merton framework (Black and Scholes 1973; Merton 1973), I value the 2024 vintage transformer slot inventory of a hypothetical operator as a deep in-the-money call option on accelerated compute deployment, written on the acceleration itself with the carrying cost of pulling the build forward as its strike, and I find around $19.7 million of value per 75 megawatt slot. I then model an Arrhenius-form insulation aging law (IEEE 2011), integrated over monthly climate normals rather than evaluated at an annual mean, under which a Miami-sited transformer consumes insulation life 2.7 times faster than one in Northern Virginia and 6.0 times faster than one in Helsinki. From these two pieces, I construct a two-name relative-value thesis (a long-biased relative-value book: long the scarce layer, funded by underweighting the crowded one, not a market-neutral hedge). One leg is long an operator with 2024 vintage transformer slots and Sun Belt exposure (because the slot optionality dominates the thermal liability). The other leg is long an operator with concentrated Nordic exposure (because the thermal asymmetry is structurally durable); that leg is a largely private universe, expressible in public markets only indirectly. The combined two-name long book represents around $380 million of value I do not find identified in public research, roughly half a percent of the market capitalisation of the closest listed comparable: a sleeve-level claim, not a re-rating call.

Keywords: Real options, Black-Scholes-Merton, transformer infrastructure, Arrhenius kinetics, IEEE C57.91, Power Usage Effectiveness, Grain-Oriented Electrical Steel, FERC interconnection, sovereign AI, geographic-thermodynamic arbitrage.

JEL Classification: G12, G13, L94, Q42, O33, R33.

0.0 Preamble

Right now, investor positioning is heavily concentrated long the silicon (NVDA, AVGO, TSM) and long on the GPU adjacent companies (CRWV, VRT, SMCI). The thinking behind this is that

the main thing slowing AI compute deployment is the supply of NVIDIA chips. That premise is wrong. The actual bottleneck is the electrical apparatus that sits between the utility grid and the rack: Generator Step-Up (GSU) transformers, primary substation transformers in the 138 / 34.5 kV class, and the auxiliary equipment that surrounds them.

Three divisions, Hitachi Energy, Siemens Energy Grid Technologies and GE Vernova Electrification, dominate OECD supply of the ≥100 MVA substation and generator step-up classes that data centers order; the three largest suppliers take roughly 45% of global power-transformer revenue on published market data, and their share of the OECD high-voltage segment specifically is materially higher (Mordor Intelligence 2025; NIAC 2024). Hitachi Energy alone has seen its order backlog grow from about $31 billion at the end of FY2023 to $57.9 billion at the end of FY2025 (Hitachi Energy 2024, 2025). Lead times for transformers above 100 megavolt-amperes have stretched from around 50 weeks in 2021 to roughly 128 to 144 weeks on average by 2026, with generator step-up and substation units quoted as far out as 210 weeks (Wood Mackenzie 2024; NIAC 2024). The market reads this as a supply story for the manufacturers, and it has bid up Hitachi Energy and Siemens Energy on that basis. What the market has not done is price the asymmetric advantage held by the *operators* who got their orders in during the 2023 and 2024 window.

This paper sets out four theses inside that mispricing. The first three find the option value, the liability, and the moat inside a hypothetical Operator X with concentrated 2024 vintage slot inventory and Sun Belt fleet exposure. The fourth thesis is the positive mirror of the second. Where Sun Belt placement creates a thermal liability under the Arrhenius aging law, Nordic placement creates a thermal asset under the same law. Operator H (the Nordic long leg) is the pair to Operator X. Combined, the two operators carry around $380 million of off-balance-sheet value I do not find identified in public research. Both operators are hypothetical throughout Section I, and every slot count and dollar figure is a stylized modeling assumption rather than any company's disclosed number. Section II names the listed vehicles that come closest to each profile and states the claim there as a claim about mechanism, not about a specific dollar mispricing at a specific company. Read the per-slot figure as the unit result and scale it yourself; that is the number the framework produces, and the aggregate is only ever sixteen times it.

Thesis I.A, The Transformer Optionality Thesis

A.1 Market Consensus

The sell-side treats secured transformer orders as prepaid CAPEX. They get written down on a straight-line basis as the deliveries come in. Equity research models the 2024 order at its historical invoice price plus a modest escalator. The gap between the 2024 strike price and the 2026 spot replacement price gets booked as a one-time cost-of-goods benefit, smoothed through gross margin. I have not found a coverage model that attributes a separate valuation to the slot itself.

A.2 Variant Perception

A secured slot is not prepaid CAPEX. It is a deep in-the-money call option on accelerated AI compute deployment, with a time-value component that GAAP statements do not capture. The

slot gives the holder the right to energize a critical electrical asset on a delivery date that competitors cannot match at any price until the supplier queue clears, which is 2028 to 2030 at the earliest. The mispricing is not in the cost of the transformer. It is in the present value of the revenue captured during the window of time when competitors are locked out.

A.3 Underlying Asset Decomposition

The underlying asset of the call option (S) is the risk-adjusted present value of the cash flow an accelerated 75 megawatt deployment block brings forward. The strike (K) is the marginal buildout CAPEX with the transformer cost excluded (the remaining incremental civil, electrical, and switchgear capital, excluding committed shell and land and tenant-owned IT equipment). The technical anchor for sizing the tenant's willingness to pay is the conversion of kilowatt-hours to revenue for an H100-equivalent training cluster:

In this equation, is the IT-load utilization factor, modeled at 0.78 for training-dominant workloads. PUE is the Power Usage Effectiveness (Green Grid 2012), set at 1.18 for a direct-to-chip liquid-cooled facility. is the marginal contracted GPU-hour price, modeled at $2.45/hr for an H100-equivalent on 2026 spot (NVIDIA 2024). 8760 is the number of hours in a year, and NGPU/MW is the accelerator-equivalents per megawatt of IT capacity, about 1,000 (1 MW divided by roughly 1 kW per node). The calculation yields around $14.2 million of gross annual revenue per megawatt, falling to around $4.2 million per megawatt in EBITDA after power, staffing, and network costs.

On business model: the $14.2 million per megawatt above is the gross compute value a fully-utilized slot unlocks for the tenant, not the revenue Operator X books. A power-secured landlord such as Digital Realty captures a fraction of it through lease economics, on the order of $1.5 to $2.5 million per megawatt. The slot's value to the landlord, and the underlying S of the option, is the scarcity rent it earns by energizing ahead of the interconnection queue: premium pre-leasing and yield uplift while capacity is short. That is why the strike K excludes GPU procurement, since the landlord does not own the GPUs, and is limited to the marginal civil and electrical build.

Three PUE figures appear in this paper and refer to different facility types, not different assumptions about the same site: 1.18 for a new direct-to-chip liquid-cooled build (used in the revenue bridge above), ~1.30 as the sell-side flat-fleet consensus, and 1.42 as the Miami air-cooled worst case (used in the thermal section). Each is applied to the matching case so the unit economics and the thermal liability are computed on a consistent, like-for-like basis.

- **What it is:** the bridge from megawatts to revenue.
- **Source:** standard data-center accounting (IT load × hours × price per GPU-hour, divided by PUE for cooling overhead).
- **Why it applies:** it converts a slot's power capacity into the cash flow it unlocks.
- **How it's used:** it sizes the tenant's willingness to pay (about $14.2M gross and $4.2M EBITDA per MW per year); S, the option's underlying, is set from the landlord's scarcity-rent capture, not this full tenant figure.

A.4 Financial Driver, Real-Options Heuristic, Black-Scholes-Merton Parameterization

Framing. Each secured 75 MW transformer slot grants the operator the right but not the obligation to deploy AI compute capacity ahead of competitors who must wait through the FERC interconnection queue. That right has the structural shape of a call option on accelerated deployment: the operator pays a fixed marginal CAPEX (the strike) and receives a stochastic future cash-flow stream (the underlying) over a window that closes when the queue clears or competitors catch up (the expiry). Following the real-options framework of Trigeorgis (1996) and Dixit and Pindyck (1994), I parameterize this as a Black-Scholes-Merton call to anchor the order of magnitude.

This is presented explicitly as a real-options heuristic, not as a defensible point estimate. The Black-Scholes preconditions of tradeable underlying, continuous delta-hedging in that underlying, log-normal price dynamics, known constant volatility, and no payout leakage do not hold for a non-traded, non-replicable transformer slot (Hull 2021). The volatility input $\sigma \approx 35\%$ used below is a proxy from H100 GPU spot-price 12-month realized volatility, applied on the working assumption that the volatility of accelerated AI-compute NPV is anchored to the volatility of the compute revenue stream. A 5-percentage-point change in σ moves the per-slot value by approximately \$0.5M; a 20% change in the underlying S moves it by approximately \$10M; the full sensitivity surface is documented in the model's `20_Robustness` tab.

A defensible interpretation of the framework is therefore: the per-slot real-options value is **order-of-magnitude \$15M to \$25M** depending on parameter choices, with a central illustrative value of **\$19.7M**. Aggregated across the sixteen secured slots and adjusted for non-tradeability, Operator X's off-balance-sheet option value is **order-of-magnitude \$240M to \$400M**, with a central illustrative value of **\$307M**. The thesis depends on the directional sign (the option value is unambiguously positive), the relative size to current market capitalization, and the asymmetry between the legs of the book, not on the precision of the central illustrative values.

A candid note on what the option model is doing here: with S well above K, the slot sits deep in-the-money, so Black-Scholes adds only about \$0.9M of time value over the plain intrinsic value (S minus the discounted strike). The headline per-slot figure is therefore driven by S, the risk-adjusted value the slot pulls forward, not by the option mechanics. I keep the Black-Scholes frame because it disciplines the inputs and bounds the time-value, but the load-bearing assumption is S, and that is where a reader should push.

One practitioner reviewer pressed precisely this input, arguing the GPU-spot proxy is not the right object. He is right, and the correction runs in the thesis's favor. The relevant volatility is that of the slot's accelerated scarcity-rent present value, which is contracted, escalator-bearing landlord income and is therefore lower than GPU-spot volatility, not higher. Across a 15 to 60 percent volatility band the per-slot value moves only from \$18.8M to \$22.6M and the 16-slot book from \$301M to \$361M, so the option is at least 83 percent intrinsic. A lower, more defensible volatility pushes the option deeper in-the-money and makes the headline more intrinsic, not less. The criticism is accepted as to the input and does not move the conclusion.

Note on S: it is the incremental, risk-adjusted value the slot pulls forward by clearing the interconnection queue early, not the full standalone NPV of the 75 MW block (which is materially larger). The option is written on the acceleration, not on the asset itself.

How S is built, and what the strike actually is. S is the present value of the rent a 75 MW slot pulls forward by energizing roughly two years ahead of queue-bound competitors. The scarcity premium a power-secured operator earns during the shortage is taken at about $0.40 million per megawatt per year, the increment over normalized lease economics rather than the full lease, and it is already risk- and probability-adjusted. Seventy-five megawatts gives $30 million a year for two years, arriving at roughly t = 1.5 and t = 2.5, discounted at the risk-free rate: 30/1.044^1.5 + 30/1.044^2.5 = $55.1 million. Earlier drafts summed the two years undiscounted and called the result a present value, which was wrong: risk-adjustment produces a certainty equivalent and time value still has to be applied, as it is to the strike inside the same formula. The inputs bracket it: at a $0.30 to $0.50 million per-megawatt premium and a 1.5 to 2.5 year head-start, S runs roughly $32 million to $86 million. Because S is the load-bearing input, that range, not the point estimate, is what the thesis rests on. The strike has to sit on the same basis. If the option is written on the acceleration, then what exercise costs is not the whole build but the carrying cost of pulling that build forward. A 75 MW fit-out excluding shell, land, the transformer itself and tenant-owned IT runs about $6.0 million per megawatt on current construction benchmarks, $450 million for the block; accelerating it by two years at the risk-free rate costs 450 × (1 − e^(−0.044 × 2)) = $37.9 million. That is K. Pairing a $61 million acceleration increment against a $42 million full build cost, as earlier drafts did, put the numerator and the denominator on different objects. The option’s T of one year is the time to the slot-delivery decision, distinct from the two-year scarcity-rent window that S values after delivery.

Three modeling caveats. Black-Scholes assumes European exercise, while the slot is closer to an American option exercisable when the queue clears; for a deep in-the-money call with no dividend-like leakage early exercise is not optimal, so the American and European values converge and Black-Scholes is a conservative approximation rather than a misuse. The sigma = 35% input is a deliberately high volatility proxy, but because the option is deep in-the-money its vega is small and the value barely moves with sigma, as the roughly $0.9 million of time value above shows. And S is itself a risk-adjusted present value, so the risk discount sits inside S rather than in a separate hurdle-rate denominator applied to a speculative cash-flow stream.

The closed-form calculation below is the central-illustrative-value derivation. I treat each 75 megawatt slot as a European call (Black and Scholes 1973):

- **What it is:** the Black-Scholes-Merton call formula, with its d1 and d2 terms.
- **Source:** Black & Scholes (1973) and Merton (1973).
- **Why it applies:** a secured slot is the right, not the obligation, to deploy compute early, which is structurally a call option (S = value pulled forward, K = build cost, T = exclusion window).
- **How it’s used:** a heuristic to size the hidden value to an order of magnitude (about $19.7M per slot), not a precise valuation, since a slot is not a tradeable, continuously-hedged asset.

The inputs for a single 75 megawatt slot are below.

Parameter	Symbol	Value	Source / Logic
Underlying, risk-adjusted PV of accelerated cash flow	S	$55.1M	PV of the scarcity rent a slot earns by energizing about two years ahead of the queue: $0.40M/MW/yr over normalized lease economics x 75 MW, at t = 1.5 and 2.5, discounted at r. See "How S is built". Order-of-magnitude input.
Strike, carrying cost of accelerating the build	K	$37.9M	What exercise costs: carry on pulling the fit-out forward two years, not the fit-out itself. Fit-out ex shell, land, transformer and tenant IT ≈ $6.0M/MW (Turner & Townsend 2025) = $450M; carry = 450 x (1 − e^(−rΔ)), r = 4.4%, Δ = 2 yr.
Risk-free rate (U.S. Treasury)	r	4.40%	Conservative; the 1-yr UST (FRED DGS1) was 3.83% on 1 June 2026

Parameter	Symbol	Value	Source / Logic
Volatility, annualized σ of AI compute capacity NPV	σ	35.0%	Realized 12-month volatility of H100 spot rental pricing across major providers. A deliberately high proxy: contracted landlord scarcity rent is less volatile than GPU spot, and a deep in-the-money option is nearly insensitive to it.
Time to expiry, slot-delivery decision	T	1.00 yr	Time to the Operator X delivery/energization decision; the two-year scarcity-rent window after delivery is embedded in S.

The calculation runs as follows:

The aggregate option value for a hypothetical Operator X with sixteen secured slots (16 × 75 MW = 1.20 GW) is:

After a single $8M liquidity and non-transferability discount (about 2.5%; the option is non-tradeable):[1]

[1] The $8M liquidity discount (about 2.5% of gross slot value) is a working assumption rather than a market-observed discount. There is no secondary market for transformer slots, so no

That $307 million option leg is not separately identified in any coverage model I have been able to review. It is not in the working capital line. It is not in deferred revenue. It is not in property, plant, and equipment. It is an off-balance-sheet real option whose value is realized as each slot gets delivered and energized (Myers 1977; Trigeorgis 1996).

On scale: against a Digital Realty-sized market capitalization (about $71B as of 31 July 2026), the $276M net is under 0.4%, a sleeve-level alpha rather than a re-rating; the same value against a mid-cap operator's enterprise value is a headline number. The edge concentrates where slot inventory is large relative to enterprise value, which is why the tradeable expressions in Section II are sized as portfolio sleeves, not single bets.

Thesis I.B, The Thermal Derating Liability

B.1 Market Consensus

Operators following a "Sun Belt strategy" (Miami, Houston, Phoenix, Atlanta) get credited with lower land cost, faster permitting, and proximity to fiber landings. The transformer fleet at these sites gets depreciated on the standard useful-life assumption of 180,000 service hours, which is about 20.5 years (IEEE 2021; NEMA 2013). Equity research treats this as climate-neutral.

B.2 Variant Perception

The aging of transformer insulation is governed by an Arrhenius-form kinetic law (Arrhenius 1889; IEEE 2011). Hot-spot temperature is the sum of three terms: ambient, the top-oil rise, and the winding-to-oil gradient. Only the last two depend on load, so at constant load ambient passes through to hot-spot one for one. That additive structure is what makes site selection a thermal decision. IEEE C57.12.00 rates a unit at a 30 °C average ambient, and the 110 °C reference hot-spot in the aging law is that ambient plus a 65 °C average winding rise and a 15 °C hot-spot differential. Miami's annual mean ambient of 25.3 °C (NOAA 2025) therefore sits below the IEEE rated ambient, not above it, and at rated load no site in this study ages faster than the laboratory reference. That is the correct absolute statement and it is not the one the thesis needs. The claim here is relative, and because the law is exponential in temperature the relative spread is large: integrating the aging factor over monthly climate normals rather than evaluating it at the annual mean, which matters because the function is convex and the mean understates it, Miami consumes insulation life 2.7 times faster than Northern Virginia and 6.0 times faster than Helsinki. In peak summer the 95th percentile Miami ambient reaches 32 to 34 °C, above the rated condition, so the site also runs with less overload headroom on the days when load peaks. The equivalent aging factor runs at 0.645 in Miami against 0.237 in Northern Virginia, an index of 2.72 on the benchmark market. The operator's replacement-capital reserve is understated on a physical basis wherever the fleet is Sun Belt weighted, and hot-spot temperature above the rated condition raises moisture, gassing and through-fault exposure during heat events (CIGRE 2015).

---

traded illiquidity spread exists; the figure is calibrated by analogy to private-equity NAV discounts on infrastructure assets with comparable resale frictions.

B.3 Cooling Oil Heat Balance and Arrhenius Aging

A 100 megavolt-ampere power transformer dissipates roughly 0.35% of its rating at rated load: about 40 to 60 kilowatts of no-load (core) loss and 250 to 350 kilowatts of load (I^2R) loss, a full-load efficiency near 99.6% (Heathcote 2007; Hitachi Energy 2025). At full load:

The cooling fluid mass flow rate required to hold the oil rise to $\Delta T = 20$ °C is found from the steady-state energy balance (Lienhard and Lienhard 2019):

For mineral oil (,):

For natural ester oil (FR3-class, ,):

- **What it is:** the coolant flow rate needed to hold the oil's temperature rise to a target.
- **Source:** the steady-state heat balance (Lienhard & Lienhard 2019): heat removed = mass flow x specific heat x temperature rise.
- **Why it applies:** it sizes the circulation duty each fluid must carry for a given temperature rise.
- **How it's used:** to compare mineral against ester fluids on circulation duty. Note that under ONAN/ONAF the oil circulates by thermosiphon, so this is a heat-balance illustration, not a pump-power calculation.
- **What it is:** the waste heat a large transformer sheds at full load.
- **Source:** transformer engineering, roughly 0.35% of rating lost to copper and core losses at rated load (Heathcote 2007).
- **Why it applies:** that heat is what cooling must remove, and heat drives insulation aging.
- **How it's used:** it feeds the cooling flow-rate calculation and links to the Arrhenius aging law below.

So ester fluids carry about 11% less mass flow for the same thermal envelope on specific heat alone. That is not an operating-cost saving. Under ONAN/ONAF cooling the liquid circulates by natural thermosiphon and there is no oil pump; the auxiliary load is radiator-fan power, which is set by the air side, and the whole cooling auxiliary budget for a 100 MVA unit is on the order of 7 to 8 kilowatts. Circulation work, where it exists in forced-oil classes, scales with volumetric flow and viscosity rather than with mass flow, and natural ester is roughly three times as viscous as mineral oil at operating temperature, so esters are a thermal penalty bought for fire safety and biodegradability. The ester case that does carry economic weight is different and stronger: at equal hot-spot temperature, paper aged in natural ester lasts roughly five to eight times longer than the same paper in mineral oil, through moisture scavenging and suppressed cellulose hydrolysis. Flow figures scale linearly with the assumed loss rate.

The Arrhenius insulation aging law (IEEE C57.91-2011, Clause 5.2; 65 °C-average-winding-rise thermally upgraded kraft, 110 °C reference hot-spot) is:

- **What it is:** the insulation aging-acceleration factor.
- **Source:** the Arrhenius reaction-rate law (1889), codified for transformers in IEEE C57.91; at B = 15,000 and a 383 K reference, aging doubles per 6.8 C of hot-spot rise.
- **Why it applies:** it turns site ambient into a relative rate of insulation life consumption, the core of the geographic case.
- **How it's used:** integrated over monthly normals it runs 0.645 at Miami against 0.107 at Helsinki, a 6.0x relative aging spread that drives the replacement-reserve and insurance deltas in the tables that follow.

Site	Annual mean ambient	Mean hot-spot at rated load		Relative aging index (NoVA = 1.00)
IEEE rated condition	30.0°C	110.0°C	1.000	4.22
Northern Virginia (Loudoun)	13.0°C	93.0°C	0.237	1.00
Dallas-Fort Worth	19.6°C	99.6°C	0.464	1.96
Miami-Dade	**25.3°C**	**105.3°C**	**0.645**	**2.72**
Phoenix	24.0°C	104.0°C	0.740	3.12

The circulation-duty sensitivity to fluid specific heat for a Miami-sited 100 MVA unit (θa = 25.3 °C) is below. It is a heat-balance comparison, not an operating-cost comparison: see the note that follows.

Fluid	(J/kg·K)	at ΔT = 20 °C (kg/s)	at ΔT = 15 °C (Miami-derated)	Mass-flow Δ vs mineral
Mineral oil (Nynas Nytro 10X)	1,860	9.41	12.55	baseline
Synthetic ester (Midel 7131)	1,950	8.97	11.97	−4.6%
Natural ester (FR3)	2,090	8.37	11.16	−11.0%
Silicone fluid (50 cSt)	1,510	11.59	15.45	+23.2%

B.4 Financial Driver, Replacement-Capital Reserve and Insurance Repricing

Expected life below combines thermal aging with the non-thermal failure mechanisms that actually retire transformers, bushings, tap changers, moisture ingress and through-faults, on a competing-risks basis with a 45-year non-thermal characteristic life (CIGRE 2015). Reporting thermal aging alone produces service lives of 30 to 270 years, which no utility fleet observes and no auditor would depreciate to. The financial object is the annual replacement-capital reserve, not book depreciation: for an operator valued on funds from operations, a depreciation schedule is added back and only the reserve, the residual value and the insurance line carry cash consequence.

Site	Expected life, competing risks (yr)	Annual replacement reserve ($9.2M unit)	NPV Δ vs. NoVA benchmark (20-yr, r = 8%)
Loudoun	29.6	$311K/yr	benchmark
DFW	22.3	$413K/yr	− $1.00M
Miami-Dade	**18.6**	**$494K/yr**	**− $1.80M**
Phoenix (peak-weighted)	17.2	$536K/yr	− $2.21M

For a hypothetical Operator X with sixteen Miami-sited 100 MVA units, the portfolio impact is:

Line Item	Value	Notes
Incremental replacement-capital reserve	$2.93M/yr	(16 × $183K Δ)
20-yr NPV at 8% discount	−$28.8M	PV annuity factor 9.818
Insurance premium uplift (catastrophic failure risk)	1.4× base (a +40% uplift)	4–7× during ≥40°C events; illustrative reserve, not an underwriting quote
Annual insurance Δ ($0.44M fleet base)	+$0.18M/yr	Base rate 0.30% of the $9.2M unit value (16 × $27.6K/yr); benchmark band 0.2% to 0.5%, derivation below
20-yr NPV insurance Δ at 8%	−$1.7M	
Total hidden thermal liability NPV	**−$30.5M**	Reserve and insurance effect on funds available for distribution; rounds to the ≈−$31M working figure

The insurance base is set at 0.30% of unit replacement value, $27.6K per unit per year, taken from inside a 0.2% to 0.5% benchmark band rather than from an underwriting quote. The band is what published pricing supports for large engineered property: onshore-wind physical-damage cover runs 0.3% to 0.4% of insured values (UNEP n.d.), the roughly $10 billion of annual data-centre premium S&P Global Ratings projects against an insurable asset base above $2 trillion

implies about 0.5% across all lines combined, not property alone (S&P Global Ratings 2026), and commercial property rates have been falling from these levels since 2024, with US property rates down about 10% year on year in early 2026 (Marsh 2026). Earlier drafts assumed $200K per unit, a 2.2% rate with no market support; v39 re-bases the line and applies the correction symmetrically to the Nordic insurance credit in Section I.D.4, which moves value between the two legs but leaves the combined book unchanged. Two caveats survive the re-basing. Utility owners of transformer fleets sometimes self-insure rather than buy cover, as rate-case testimony documents for Texas transmission fleets (Entergy Texas 2013); in that case this line is a self-insurance reserve rather than a premium, with the same cash consequence for funds available for distribution. And Miami carries a windstorm-catastrophe loading on top of any thermal effect, with Florida coastal commercial property priced materially above national levels (Federal Reserve 2025); it is excluded here because it prices hurricane exposure, not insulation kinetics, and including it would overstate the thermal claim.

This is a negative mispricing, and the liability is not in the operator's investor relations deck. It offsets about 10% of the Thesis I.A optionality value, but it is site-mix dependent. An operator with a fleet weighted more than 50% to Loudoun gets a positive kickback (extended life), while an operator weighted more than 40% to the Sun Belt is structurally short the insulation kinetics.

This is an economic liability, not a claim that the GAAP books are wrong. Statutory depreciation follows MACRS or straight-line schedules and does not track localized insulation aging; the mispricing is the unreserved replacement capital, higher insurance, and lower residual value that faster physical aging imposes, none of which the straight-line schedule captures.

Thesis I.C, The Regulatory and Manufacturing Choke as a Compounding Moat

C.1 Market Consensus

FERC interconnection queues and NEPA reviews show up in the financial press as headwinds. Bureaucratic friction that will eventually clear once Congress passes some form of permitting reform (FERC 2023; U.S. Congress 1969). The sell-side base case assumes a reversion to mean in interconnection-to-energization timelines. Right now the national median from interconnection request to commercial operation runs beyond five years for projects energized in 2025, against 4.5 years for the 2024 cohort (Rand et al. 2025). The base case targets 2.5 years by 2028.

C.2 Variant Perception

The interconnection queue and the transformer manufacturing queue are structurally coupled. They reinforce each other. Even if FERC cuts review timelines by half, that does not accelerate energization for any operator who has not already secured a transformer slot. The manufacturing oligopoly (Hitachi, Siemens Energy, GE Vernova) is constrained on three fronts at once, grain-oriented electrical steel, winding capacity and certified craft labour, and the Department of Energy reports US large-power-transformer plants running near 40% capacity utilisation, which points at labour and tooling rather than steel alone as the near-term binding input (NIAC 2024; DOE 2024; USGS 2025). The moat is therefore compound. A 2024-secured slot is protected by a

regulatory wall (FERC), a metallurgical wall (GOES), and a labor wall (the high-voltage winding craft has a seven-year apprenticeship cycle). The sell-side treats these as independent risks. They are not: within the manufacturing cluster the average pairwise correlation runs about 0.70, with the strongest pair (GOES and labor) near 0.83 (Appendix B).

C.3 Grain-Oriented Electrical Steel as a Magnetic Flux Constraint

The cores of high-voltage transformers require Grain-Oriented Electrical Steel with specific magnetic permeability characteristics along the rolling direction (Bozorth 1993; Fitzgerald, Kingsley, and Umans 2014). The core loss per kilogram is given by the modified Steinmetz equation (Steinmetz 1892):

In this equation and are the hysteresis and eddy-current coefficients, is 60 Hz, is the peak flux density (typically 1.7 Tesla for HiB-grade GOES), and is around 1.6. A 100 megavolt-ampere transformer needs roughly 28 tonnes of GOES in its core (Nippon Steel 2024). Annual OECD production of high-voltage-transformer-grade GOES runs on the order of 480 kilotonnes, against roughly 5,000 large power transformers (≥100 MVA) built globally each year at an average well above 100 MVA, plus medium-power units, shunt reactors and 10 to 15% core-cutting scrap that draw on the same grades. GOES is therefore not a hard volume ceiling on its own; it is a cost and qualification constraint, more than 20% of transformer cost, on a grade base that cannot be requalified quickly (IEA 2024; DOE 2024). AI data-center transformers absorb an estimated 35 to 40% of the annual growth in GOES demand in 2026 (central case about 36%, range 20 to 50%; see Appendix B). Growth share, not stock share, is the defensible claim.

The moat compounds because every layer is bottlenecked independently.

Constraint Layer	Resolution Lead Time	Counterparty Concentration	Substitutability
FERC interconnection (PJM)	>5 yr median (national)	Single ISO	None
NEPA Environmental Assessment	1.8 yr median	Federal agency	Categorical exclusion (limited)
HV transformer manufacture	130–210 wk	3 OEMs (Hitachi/Siemens/GE Vernova)	None (GOES-bound)
GOES (HiB-grade) supply	Mill-bound, 12+ mo	About 6 HiB-capable mills (Nippon Steel, JFE, POSCO, thyssenkrupp, Baowu, Cleveland-Cliffs)	None
HV winding craft labor	7-yr apprenticeship	<2,000 certified globally	None

Measured on year-over-year changes over 2018-2025 (Appendix B), the three manufacturing-supply constraints - transformer lead time, GOES tightness, and high-voltage winding labor - move together rather than independently:

Transformer-GOES 0.52, transformer-labor 0.74, GOES-labor 0.83, for an average pairwise correlation near 0.70.

A single component explains about 80% of that cluster's variance (Jolliffe 2002), so the three behave as one underlying supply factor. Treating them as independent, as the sell-side does, understates the 95th-percentile energization delay by roughly four months, not several years. Interconnection-queue and NEPA timelines are driven by separate dynamics - NEPA review times actually fell over the period - so they are not part of this latent factor, which is why the moat is sized on the manufacturing cluster rather than on all five constraints.

C.4 Financial Driver, The Moat Premium

The 2024-secured slot inventory therefore carries a compounded scarcity premium during the 2026 to 2030 window. This premium is *additive* to the Black-Scholes time value of Thesis I.A, because Black-Scholes assumes the underlying is spot-tradeable. The Operator X slot inventory is among the limited deliverable inventory during the exclusion window. Earlier drafts carried this as a separate additive line of $24 million, a 6.9% uplift attributed to correlated-constraint risk. That line is retired here. There is no clean route from a pairwise correlation to a percentage value uplift, the effect it describes is dispersion in the queue-clearing date and therefore already belongs in the volatility input rather than in a second additive term, and correlated constraints can clear together faster than the median as readily as slower, so booking only the slow tail is asymmetric. The correlated-constraint finding stands as a statement about the shape of the exclusion window; it is not separately capitalised.

Adjustment	Value (Operator X aggregate)	Rationale
Base acceleration-option value	$315M	Section A.4
(−) Liquidity / non-tradeability discount	−$8M	2.5% haircut for slot non-fungibility
(−) Thermal derating liability (Section B.4)	−$31M	Reserve, insurance and residual-value effect; site-mix dependent
Net hidden value (Operator X, off-balance-sheet)	**+$276M**	Conservative band: $260M – $320M

Net of the Miami thermal liability, the Operator X bridge totals $276 million; the $307 million liquidity-adjusted option leg is its largest component. Against a Digital Realty-sized capitalisation that is under 0.4%, a sleeve-level claim, and the paper says so rather than calling it a re-rating.

Modeling note. The constraint-correlation result, the principal-component summary, the energization-delay simulation, and the GOES demand-share figure are built from sourced 2018-

2025 series and reproduced in Appendix B and the accompanying workbook. Two figures from earlier drafts are corrected there: the independence assumption understates the 95th-percentile delay by about four months rather than 3.4 years, and AI absorbs roughly 35 to 40% of GOES demand growth rather than 42% of marginal supply. The principal-component share is reported here at about 80%; a sensitivity check records 76% under the primary construction and 60 to 85% across alternative constructions, and the 76% figure governs. The sub-2,000 certified high-voltage winders is a sourced estimate.

Thesis I.D, The Thermodynamic Geography Arbitrage (Nordic Long, Sun Belt Liability)

D.1 Market Consensus

Sell-side equity research treats data-center site selection as a function of four things: fiber proximity to demand centers, land cost, tax incentives, and regional wholesale power price (Cushman & Wakefield 2025; JLL 2025). Climate typically shows up in operator pro formas only as a baseline PUE assumption, usually a flat 1.30 across the U.S. fleet (Uptime Institute 2024). It does not flow through to differentiated transformer depreciation, insurance, or NPV. The market treats Miami and Helsinki as climate-neutral.

This is the same mistake identified in Thesis I.B, viewed from the other side. If Sun Belt placement is a liability the consensus has missed, Nordic placement is an asset the consensus has equally missed.

D.2 Variant Perception

The Arrhenius law that consumes insulation life at Miami's 25.3 °C annual mean ambient preserves it at Helsinki's 6.1 °C annual mean (FMI 2025). The same equation that gave in Thesis I.B gives at Helsinki and at Luleå inside the Arctic Circle (SMHI 2025). Thermal aging alone would imply service lives of two centuries at those latitudes, which is why the financial layer below is built on a competing-risks life that caps the benefit at the point where non-thermal failure mechanisms bind: expected life runs 18.6 years at Miami against 36.5 at Helsinki and 38.7 at Luleå, roughly a doubling rather than an unbounded extension. On top of that, the free-cooling regime (operation without compressor-based chillers) covers 85% to 95% of annual hours in the Nordic latitudes versus less than 10% in Miami (ASHRAE 2021). This collapses realized PUE from 1.42 to between 1.07 and 1.10.

These are not small improvements. They are durable advantages that cannot be replicated by Sun Belt operators without physically moving. The proof is in where the hyperscalers have actually built. Google has operated a major facility in Hamina, Finland since 2011 (site acquired 2009), using Baltic seawater for cooling and reporting a PUE between 1.10 and 1.12 (Google 2024). Meta operates the Luleå Node Pole hyperscale campus inside the Arctic Circle (Meta Platforms 2023). Nebius, formerly Yandex's international business, operates at Mäntsälä, Finland, selling waste heat into the municipal district-heating grid (Nebius 2024). Microsoft is commissioning a major facility in the Espoo region of Finland in 2025 (Microsoft 2023). These four companies have built hyperscale capacity in the Nordic region based on internal infrastructure economics. They have not clustered at high latitude by accident. They are validating math I do not find consistently reflected in public models of mid-tier operators.

The variant perception is therefore not the existence of the Nordic advantage (which the hyperscalers have already demonstrated through where they actually put their money). It is the absence of that advantage from the sell-side models of publicly-traded operators with Nordic exposure.

D.3 Nordic Recalculation of

Using the same IEEE C57.91 form from Thesis I.B.3 (IEEE 2011):

Hot-spot temperature is calculated under the same ONAN / ONAF cooling convention as the Sun Belt sites in Thesis I.B.

Site	(°C)	(°C)		Expected life, competing risks (yr)	Free-Cooling Hours / Year
Miami-Dade	25.3	105.3	0.645	18.6	~500 (~5.7%)
DFW	19.6	99.6	0.464	22.3	~3,800 (~43%)
Loudoun (NoVA)	13.0	93.0	0.237	29.6	~5,200 (~59%)
Helsinki, FI	**6.1**	**86.1**	**0.107**	**36.5**	**~7,800 (~89%)**
Stockholm, SE	7.6	87.6	0.120	35.6	~7,500 (~86%)
Luleå, SE (Meta Node Pole)	**1.8**	**81.8**	**0.074**	**38.7**	**~8,400 (~96%)**
Reykjavík, IS	5.5	85.5	0.076	38.6	~8,000 (~91%)

A facility runs without compressor-based mechanical cooling whenever the ambient wet-bulb temperature sits below the chilled-water setpoint less the cooling-tower approach (about 5 to 6 K) and, for an indirect economiser, a further plate-exchanger approach of 2 to 3 K. The hours tabulated above are computed against that net threshold, not against the raw setpoint. Direct-to-chip facility water is trending warmer, toward 27 to 32 °C in current deployments rather than the legacy 18 to 24 °C chilled-water band (ASHRAE 2022), which raises free-cooling hours everywhere and compresses the geographic spread; the tabulated hours are held at the conservative legacy threshold, so the Nordic advantage shown here is the smaller of the two readings. The fraction of operating hours qualifying for free cooling drives the realized PUE:

In this equation, is the fraction of operating hours in the free-cooling regime, (pump and fan parasitics only), and (compressor-based chiller plant active).

- **What it is:** realized Power Usage Effectiveness as a blend of free-cooling and chiller hours.
- **Source:** the Green Grid PUE metric, weighted by the share of the year a site cools without chillers.
- **Why it applies:** cooler climates run more free-cooling hours, lowering energy cost.
- **How it's used:** it produces the 1.10 (Helsinki) vs 1.42 (Miami) PUE gap behind the Nordic OPEX advantage.

One caution on scope. The Green Grid metric puts everything inside the utility meter in the numerator, including transformer and switchgear losses and their cooling auxiliaries, so transformer-side auxiliary energy is already counted in PUE and must not be added to it. Earlier drafts claimed the two were additive; they are not. The cooling-energy line below is a PUE delta and is taken as the whole of the climate-driven auxiliary saving.

Site		Realized PUE	Cooling Energy vs. Miami
Miami-Dade	0.057	1.42	baseline (100%)
Loudoun	0.594	1.21	−50%
Helsinki	**0.890**	**1.10**	**−76%**
Luleå	**0.959**	**1.07**	**−83%**

Circulation duty also falls in the Nordics. Recalculating the Thesis I.B.3 heat balance for a Helsinki-sited 100 MVA unit under ΔT = 30 °C, achievable given the much colder heat sink:

That is a 33% reduction in circulation duty against the Miami case (kg/s at °C). Under ONAN/ONAF this is not an operating-cost saving, because the liquid circulates by thermosiphon and the whole cooling auxiliary budget for a 100 MVA unit is on the order of 7 to 8 kilowatts. Earlier drafts booked a $40 thousand per unit per year pump saving here; that line has been removed from the Operator H bridge.

D.4 Financial Driver, Operator H Construction and the Two-Name Book

The per-transformer NPV uplift for a Helsinki-sited 100 MVA unit, measured against the Northern Virginia benchmark rather than the IEEE laboratory reference, is below.

Component	Annual Δ ($K)	NPV Δ ($M)	Notes
Replacement reserve saved (vs. $311K NoVA benchmark)	−58	+0.57	$9.2M unit, expected life 36.5 yr → $252K/yr

Component	Annual Δ ($K)	NPV Δ ($M)	Notes
Insurance premium reduction	−11	+0.11	−40% vs. the $27.6K/unit/yr base set in Section I.B.4 (benchmark-band rate, not an underwriting quote)
Free-cooling chilled-water OPEX	−620	+6.09	76% cooling-energy delta on a conservative 7 MW IT per-unit allocation; back-solves to a delivered power price near $32/MWh; fuller utilization scales this line up
Total per-unit NPV uplift	**−689**	**+6.77**	Reserve, insurance and operating-cost effect

For a hypothetical Operator H with sixteen Helsinki / Luleå-weighted 100 MVA units:

Line Item	Value ($M)	Notes
Per-unit NPV uplift	+6.77	Above
Aggregate, 16 units	+108.3	Gross
(−) Latency penalty on EBITDA (inference workload, ~12%)	−22.0	Helsinki ↔ NYC = 80–120 ms
(−) Currency / repatriation friction (EUR/SEK, 4%)	−3.9	FX / cross-border risk reserve; 4% applied to the roughly $97M cooling-OPEX uplift
(+) Sovereign-AI / ESG premium (low-carbon EU-siting)	+18.0	EU-sited supply at 80 gCO_2/kWh vs. roughly 375 ERCOT (operational); assumed premium, not derived from the D.6 WACC spread
(+) District-heating waste-heat monetization (Nebius Mäntsälä precedent)	+4.0	Around $250K/MW/yr
Net hidden value, Operator H (off-balance-sheet)	**+104**	Conservative band: $95M – $125M

The Operator H bridge totals $104 million (conservative band $95 to $125 million). Free cooling, not insulation aging, is the dominant line; that is the honest ranking and it is the one that survives scrutiny, because the PUE delta is directly observable in hyperscaler disclosures while the aging benefit is bounded by non-thermal failure modes.

The two-name book is constructed below. Both legs are long; the Sun Belt thermal drag is a liability embedded inside Operator X, not a separate equity short. Built against the same risk-free and market-beta backdrop, the Arrhenius and free-cooling differential is left as the dominant differentiated alpha source. A literal long/short pair (long grid, short cooling) is set out in the companion positioning note.

Leg	Operator	Value Source	Hidden NPV ($M)
Long	Operator H (Helsinki + Luleå, 16 × 100 MVA)	Thesis I.D thermal asset	+104
Long	Operator X - Black-Scholes option value (16 slots)	Thesis I.A optionality	+315
(−)	Liquidity / non-tradeability discount	Thesis I.A.4	−8
Liability	(Implicit, embedded in Operator X)	Thesis I.B Miami thermal liability	−31
Net combined book mispricing (≈$380M)	n/a	n/a	+380

Together, the pair adds around $104 million of Operator H value to the $276 million Operator X net, for a combined book of around $380 million across the two names. That lands within 5% of the probability-weighted figure reached in a separate bear-case pass over the same inputs, which is a consistency check rather than an independent one.

D.5 Hyperscaler Revealed Preference

The Thesis I.D framework is not speculative. The most sophisticated infrastructure operators have already committed to it.

Operator	Facility	Latitude	Operational Since	Published PUE	Cooling Method
Google	Hamina, Finland	60.6° N	2011	1.10–1.12	Baltic seawater direct cooling
Meta / Facebook	Luleå “Node Pole,” Sweden	65.6° N	2013	1.04–1.07	Outside-air free cooling

Operator	Facility	Latitude	Operational Since	Published PUE	Cooling Method
Nebius (ex-Yandex)	Mäntsälä, Finland	60.6° N	2015	1.07	Free cooling + waste-heat to district heating
Microsoft	Espoo region, Finland (announced 2022)	60.2° N	2025 (commissioning)	1.10 (target)	Liquid + free cooling
Equinix	Helsinki HE7, HE9	60.2° N	2018, 2023	1.15–1.18	Outside air
Hetzner	Helsinki & Tuusula, Finland	60.2° N	2010, 2020	1.12	Free cooling
AWS	Eskilstuna, Sweden (eu-north-1)	59.4° N	2018	1.10 (estimated)	Outside air

The revealed preference is clear (Google 2024; Meta Platforms 2023; AWS 2024). The hyperscalers, whose site selection reflects detailed internal infrastructure economics, have collectively committed billions to Nordic capacity, and their disclosed PUE figures are consistent with the mechanism described here. I do not find that math reflected in public models of mid-tier operators with Nordic exposure.

D.6 Sovereign AI and the Geopolitical Overlay

Finland joined NATO in April 2023, and Sweden in March 2024. Both are EU member states. The carbon intensity of the Finnish grid is around 80 gCO_2/kWh (mostly nuclear, hydro, and wind; Fingrid 2024). The Swedish grid is around 40 gCO_2/kWh on a lifecycle basis and under 20 gCO_2/kWh operationally (hydro and nuclear; Svenska kraftnät 2024). Both sit far below the roughly 375 gCO_2/kWh ERCOT operational baseline (ERCOT 2024; EIA 2025). The Nordic figures are lifecycle and the ERCOT figure operational; the ordering holds on either basis. This differential matters in four ways.

To begin, the EU Carbon Border Adjustment Mechanism (CBAM) phases in from 2026 to 2034 (European Commission 2023). CBAM applies to imported physical goods (steel, aluminium, cement, fertiliser, electricity, hydrogen), not to digital or AI services, so it is not a direct tariff on compute. Its relevance is indirect: it raises the embedded-carbon cost of high-emission grid electricity, which favours low-carbon Nordic grids at the margin; finished IT hardware is not a CBAM good, though an extension to steel- and aluminium-intensive downstream products is proposed and pending (COM(2025) 989). Hyperscaler enterprise procurement is increasingly driven by RE100 and 24/7 carbon-free energy commitments (RE100 2024; Google 2023). Microsoft, Google, Amazon, and Meta have all published these targets, and Nordic grids meet them natively. Beyond that, EU sovereign-AI mandates (EuroHPC, Gaia-X) prefer EU-

jurisdiction compute (EuroHPC 2024; Gaia-X 2024). Lastly, post-2024 NATO defense AI procurement frameworks favor NATO-jurisdiction infrastructure for defense workloads, a dimension that is becoming more important as alliance-level AI defense initiatives expand.

The geopolitical premium attached to Nordic infrastructure is not a soft factor. It is a quantifiable component of the cost-of-equity spread between Nordic-jurisdiction and Gulf or non-aligned-jurisdiction operators (PIF and Mubadala-backed projects face EU-jurisdiction discounting in particular). A 50 to 80 basis point WACC reduction on Nordic infrastructure cash flows is defensible against Gulf or non-aligned comparators (Damodaran 2024). Against the US benchmark this book actually trades, the spread is thinner, and the sovereign-siting line in Section I.D.4 is carried as an assumed premium rather than a figure derived from this spread.

What ultimately forces the market to price this gap is verified comparative evidence arriving in front of a decision-maker with the power to compel a response, and the speed of the repricing depends on which route the evidence travels. Lyon and Maxwell (2011) model greenwash as selective disclosure disciplined by the threat of an activist audit: firms disclose their true positives, withhold their negatives, and an audit exposes what was withheld. Two features of that model do not carry over here, and the differences are load-bearing rather than cosmetic. First, the model's auditor is an activist whose audit arrives with a probability the firm takes as given; the model assigns the auditor no objective of its own. This paper substitutes the rival operator as the auditing party, since in hyperscaler infrastructure the counter-evidence, grid intensity, PUE and free-cooling hours, sits in the rival's hands in audited form. But a rival is a strategic actor: where an NGO audits and publicizes, a rival can also petition a regulator or litigate, it selects among channels by private return net of cost, and that selection makes the audit probability endogenous, and possibly zero. A rival's accusation also carries an interest discount an activist's does not, and invites symmetric scrutiny of the rival's own carbon marketing. Second, the model's own comparative statics warn that audit pressure can produce silence rather than candor, pushing some firms to disclose less; here that outcome is a Sun Belt operator quietly withdrawing its carbon claims, an event with no exposure moment and no repricing attached. What this paper borrows from Lyon and Maxwell is therefore the audit-threat framing, not the model's results, and what the framing buys is a taxonomy of channels, not a timetable. The repricing clock is not set by the regulatory reporting calendar; it is set by whether a rival acts at all, by the enforcement channel it chooses, and by the decision-maker who sits at the end of that channel.

Four channels are available. They differ in decision-maker and in binding force, and their defined clocks run from months to more than half a decade, with one channel carrying no defined clock at all.

Channel	Decision-maker	Plausible clock	State of play, August 2026
Investor disclosure: the rival publishes its audited KPIs against the laggard’s claims	Investors and allocators	Months to execute; repricing not assured	Cheapest and weakest: nothing compels a restatement, the source carries an interest discount, and the effect requires investors to treat the differential as material. The rival’s own KPIs are audited under the EED scheme, but the laggard’s facility-level figures stay confidential, so the comparison the market needs is the piece nothing compels into the open. On its own this channel may never bind.
Regulatory mandate, EU: the rival petitions for facility-level publication and rides the mandatory schedule	European Commission and national regulators	1 to 3 years, EU jurisdiction only, contingent on pending acts	The EED database (Commission Delegated Regulation 2024/1364) has collected PUE, renewable-energy factor and waste-heat reuse from every EU facility above 500 kW since September 2024, but facility-level KPIs are confidential by Article 5(5); the first aggregated release (April 2026) covered 776 facilities and 6.4 GW at a weighted average PUE of 1.36, on partial member-state coverage (European Commission 2026); a draft second-phase rating scheme with electronic labels is pending adoption; CBAM certificate sales begin February 2027 for 2026 imports (Regulation (EU) 2025/2083); CSRD wave-two reports arrive for FY2027 under the narrowed scope of Directive (EU) 2026/470.

Channel	Decision-maker	Plausible clock	State of play, August 2026
FTC enforcement, US: the rival files a Section 5 complaint alleging deceptive carbon claims under the Green Guides	The Commission	Roughly 3 to 6 years to a final order, on top of a nonpublic investigation before any complaint issues	The Green Guides remain the 2012 text (16 CFR Part 260; Federal Trade Commission 2012) and the December 2022 review is dormant; the Commission's environmental docket lists no Section 5 case since 2022 (as of August 2026); the closest precedent, ECM BioFilms (FTC Docket No. 9358, affirmed 851 F.3d 599 (6th Cir. 2017)), ran two years from administrative complaint to final order and 3.4 to appellate affirmance.
Private litigation, US: a Lanham Act §43(a) false-advertising suit, with Sherman Act §2 as the slower variant	A federal court	Lanham: months to a preliminary injunction, 1 to 3 years to judgment. Antitrust: 3 to 5 years or more	Section 43(a) gives a competitor standing (Lexmark v. Static Control, 572 U.S. 118 (2014)) and, since 2020, a rebuttable presumption of irreparable harm at the injunction stage; false scientific claims are actionable between competitors (Eastman Chemical v. PlastiPure, 775 F.3d 230 (5th Cir. 2014)); no carbon-claims case between operators is yet on record, and a first comparative carbon claim would be an establishment claim needing expert record development, so the fast end of the Lanham clock is optimistic. The §2 route must defeat the presumption that false advertising has a de minimis effect on competition (Duty Free Americas v. Estée Lauder, 797 F.3d 1248 (11th Cir. 2015)).

Three of the four clocks are drawn from published records; the EU clock is a forecast contingent on adoption of a pending act, and it is labelled as such. Trademark practice shows a median of roughly one month from filing to preliminary injunction where one is granted, and the same report notes that false-advertising claims take longer to reach relief (Lex Machina 2016); private antitrust cases reach summary judgment at a median of about 2.5 years and trial at about 3 (Lex Machina 2019); the FTC's own environmental docket supplies the ECM BioFilms timeline. The ranking that emerges is uncomfortable for any disclosure-calendar story. The fastest binding channel has no carbon-claims precedent between operators; the mandated channel publishes only

aggregates until a pending act changes that; the agency channel is currently dormant; and the channel most naturally available to a data-rich rival, handing verified comparatives to investors, is the one with no compulsion behind it.

The paper cannot determine which channel a rival will select, and it does not model the choice. Selection turns on unobservables: litigation appetite, counsel's judgment on whether specific claims are falsifiable statements of fact rather than puffery, the rival's own exposure to symmetric scrutiny, and the regulatory temperature in the jurisdiction where recognition matters. A rival may also exercise none of them, since every Nordic operator that markets its own carbon claims has something to lose from a claims war; on this section's own private-return logic, silence or the slow channels are the modal outcomes, which is exactly why timing is the binding constraint. One further dependence is owned rather than hidden: any channel that fires reprices the Nordic leg by exposing Sun Belt carbon claims, which would simultaneously force recognition of the thermal liability the Operator X bridge already carries at full value, so the two legs are correlated under this catalyst rather than independent. The timing of the Nordic repricing is therefore an open dependency of this leg, not a resolved variable, and the honest statement of the catalyst is conditional: on the regulatory or Lanham channel the repricing falls inside the Wealth Management (1 to 3 year) window; on the FTC or antitrust channel it slips beyond it; on investor disclosure alone it may not arrive at all. The conviction read splits accordingly. On the mechanism, Medium: the differential is physically real, mandatory reporting keeps accumulating the evidence, and the channels exist. On timing inside the Wealth Management window, Low, until either a rival commits to a channel or the second-phase rating scheme is adopted and converts a confidential database into published comparability; that adoption is the one named, scheduled event the clock genuinely turns on. It is the mechanism and the channel, not the magnitude, that the long leg is exposed to.

Principal Objection, Behind-the-Meter Generation and the SMR Question

The strongest argument against this entire book is that the bottleneck moves. If hyperscalers stop waiting for the grid and generate power on-site, gas turbines now and small modular reactors (SMRs) later, they bypass the FERC interconnection queue (Section I.A) and, the objection runs, the transformer oligopoly that anchors the moat (Section I.C). The demand side makes this credible: AI load is already straining regional supply, and operators are signing on-site power deals specifically to escape the queue. The argument deserves a direct answer, not a footnote.

Three things hold the thesis together.

First, timing. Commercial SMRs are a 2030s event, not a 2026-2028 one. NuScale holds the only NRC-certified design (DOE 2025); the first-of-a-kind builds - Google-Kairos (first reactor targeted 2030), TerraPower in Wyoming (~2031), Oklo on an aggressive 2027-2028 path - are pilot scale and schedule-fragile (Nuclear Innovation Alliance 2025). Nothing in the SMR pipeline touches the Trading (3-12 month) or Wealth Management (1-3 year) horizons. It is a Long-Term (5 year+) consideration, and it is priced as one.

Second, what is actually being deployed behind the meter today is gas, not nuclear, and gas does not escape the equipment queue - it joins it. The on-site projects now in build (Crusoe; the Meta-Williams campus at New Albany, Ohio, approved June 2025) run on industrial gas turbines and

reciprocating engines, while the heavy-frame units hyperscalers want at scale carry lead times of roughly three to five years; GE Vernova, Siemens Energy, and Mitsubishi each report backlogs into the early 2030s, with GE Vernova's data-center electrification orders in Q1 2026 alone exceeding its full-year 2025 total (GE Vernova 2026). Behind-the-meter relocates the chokepoint from the interconnection queue to the turbine and switchgear queue - the same suppliers, in several cases the same factories.

A first-quarter 2026 update sharpens this. GE Vernova’s gas-turbine backlog and slot reservations reached about 100 gigawatts, up 17 gigawatts in a single quarter, and management expects reservations to be sold out through 2030 by year-end; Siemens Energy closed FY2025 with a record backlog of 138 billion euros (reported November 2025), since grown to about 154 billion at the second quarter of FY2026 (GE Vernova 2026; Siemens Energy 2026). The nuclear commitments are now explicit and confirm the timing rather than break it: Meta contracted for up to 6.6 gigawatts across Vistra (about 2.6 GW: 2.2 GW of existing operating plants plus uprates), TerraPower (up to 2.8 GW) and Oklo (1.2 GW), Google for up to 500 megawatts from Kairos targeting a first reactor around 2030, and NuScale’s uprated 77 megawatt design, which received NRC Standard Design Approval in May 2025 (the 50 MWe variant holds the design certification), anchors a non-binding Tennessee Valley Authority and ENTRA1 agreement for up to 6 gigawatts. No Western small modular reactor is in commercial operation as of mid-2026. The intent is real and the delivery is 2030-plus, which is how the long-dated leg is priced.

Third, and most durable: no generation source reaches a server without a step-up transformer and medium-voltage switchgear. A behind-the-meter reactor or turbine still needs a generator step-up (GSU) transformer, switchgear, and protection, supplied by the same three-name oligopoly (Hitachi Energy, Siemens Energy, GE Vernova) that anchors Section I.C. On-site generation changes who sits in the FERC queue. It does not change who is short transformers. The grid-equipment long leg is robust to the bypass precisely because the bypass runs through the same equipment.

Where the objection does bite: if SMR timelines compress materially - HALEU fuel supply resolves, the NRC moves to fleet-style licensing, and first-of-a-kind costs fall faster than the consensus 2030s schedule - the structural (5 year+) leg of the Nordic thesis weakens, because cheap modular baseload sited anywhere erodes the geographic-thermodynamic arbitrage of Section I.D. That is the honest tail, and the reason the long-dated leg is held as a thesis rather than a high-conviction position. It is not a reason to fade the 1-3 year transformer-scarcity call. The near-term shortage is mechanical, the bypass is not yet built, and when it is built it will be built out of the same scarce iron.

Principal Objection II, Demand-Side Flexibility and the Schedulable Load

Scandizzo pressed the demand side: if compute can be scheduled, shifted in time, and moved between sites, a site-level power constraint dissolves before it binds, and the interconnection wait at the center of this paper is softer than it looks. It is the most serious version of the objection because it does not dispute the shortage of iron; it argues the shortage can be routed around in software. Flexibility is real and large, and it changes when a data center draws power, not

whether the site needs a step-up transformer and a completed interconnection to exist. It shaves the grid's peak. It does not deliver the transformer.

The objection has the best evidence of any counterargument here. Duke's Rethinking Load Growth (Norris et al. 2025) models large flexible loads across the 22 balancing authorities serving about 95 percent of US demand and finds the existing system could absorb roughly 76 to 126 gigawatts of new load if those loads accept curtailment in under 1 percent of annual hours. The mechanism it names, curtailment-enabled headroom, is the right one: a load that will briefly stand down lets the grid run existing capacity harder rather than building for a peak that occurs a few hours a year.

The number does not reach the constraint this paper is about, and the study says so. Its authors call it a first-order estimate, state that they did not model transmission constraints, and note it is unclear whether data centers can express the flexibility assumed. Curtailment-enabled headroom is measured on capacity that is already interconnected; it presupposes the step-up transformer and the energized point of connection. The thesis is about getting the connection in the first place. Flexibility raises the utilization of transformers that already exist. It does not manufacture the ones that do not, nor shorten a transformer lead time that runs years (FERC 2023; Rand et al. 2025).

Three points fix why the routing-around fails on the horizon that matters. First, the flexibility lives in the wrong half of the load: frontier training is synchronous and latency-bound, and a 200-plus megawatt cluster cannot be throttled deeply without stranding the most expensive capital expenditure in technology, nor sharded across distant nodes without destroying interconnect bandwidth, so the schedulable load is mostly inference and batch, the least transformer-hungry per dollar of value. Second, spatial shifting conserves transformer demand rather than removing it, because moving a job to a looser balancing area moves its step-up requirement with it, and the aggregate count of units the system must order is invariant to where the compute runs (Knittel, Senga, and Wang 2025). Third, the horizons do not line up: flexibility bites peak-capacity adequacy on the 2026 to 2030 grid, not the equipment lead time and interconnection wait that are load-bearing here.

The honest concession is the SMR concession again, bounded the same way. On the Long-Term (5 year+) horizon, if curtailment-enabled headroom, behind-the-meter generation, and mature workload-shifting compound, together they can defer a slice of marginal interconnection demand and soften the structural leg, which is why the Nordic long position is held as a thesis rather than a high-conviction call. It is not a reason to fade the near-term transformer call, because every mechanism in that tail still terminates at a step-up transformer. The conviction read: demand-side flexibility is a Medium-conviction mitigant to grid-peak adequacy and a Low-conviction threat to the transformer thesis on the Trading (3 to 12 month) and Wealth Management (1 to 3 year) horizons, rising to a bounded structural risk on the Long-Term leg. The names in the book do not change.

Table, Where demand flexibility acts on the thesis

Channel	What it changes	What it does not change	Net effect on the transformer call
Temporal shifting (curtailment-enabled headroom, ~100 GW; Norris et al. 2025)	Raises utilization of already-interconnected capacity	The interconnection wait and step-up transformer for any new site; transmission constraints (unmodeled)	Neutral near-term; small deferral on the 5 yr+ leg
Spatial shifting between balancing areas	Where the compute runs	Aggregate count of step-up transformers the system must order	Neutral; reallocates the queue, does not shorten it
Inference / batch interruptibility	Smooths variable, latency-tolerant load	The synchronous 200 MW+ training cluster that sets marginal rent	Neutral; flexibility sits in the least transformer-intensive load
Flexibility as a cost/emissions tool (Knittel et al. 2025)	System cost down; emissions up in carbon-heavy regions	Whether new firm capacity can be interconnected at all	Confirms flexibility is a grid-economics lever, not a bypass

Source: author's synthesis; anchors Norris et al. 2025, Knittel, Senga, and Wang 2025, FERC 2023, Rand et al. 2025.

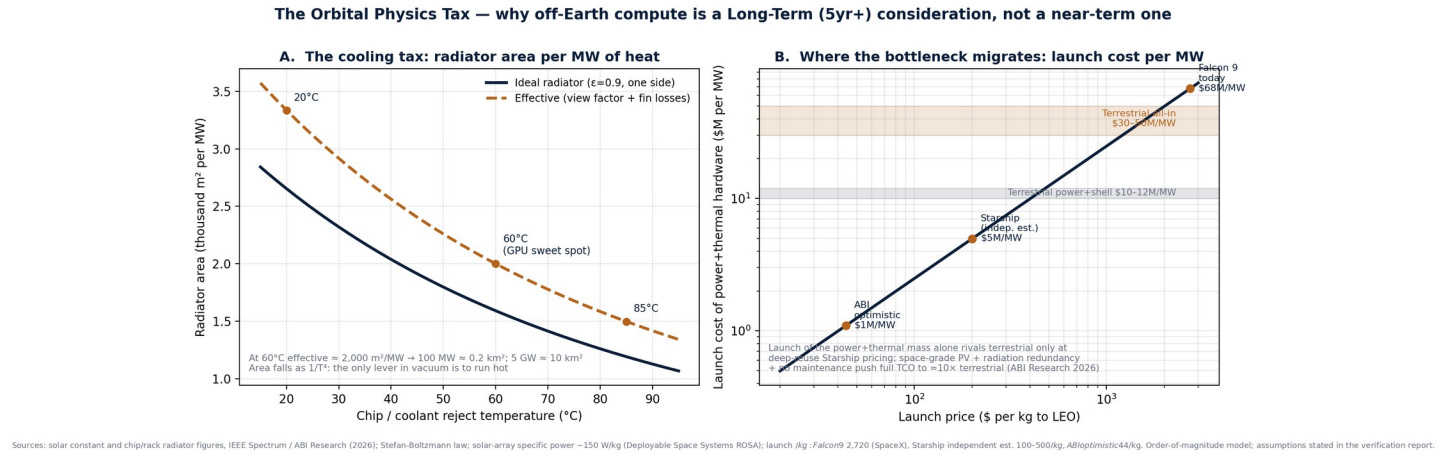
Principal Objection III, The Orbital Escape Hatch and the Terminal Migration

Section I’s behind-the-meter objection asked whether the operator can escape the grid by generating power on site. This section asks the limit case of the same question: whether the operator can escape the planet. By the middle of 2026 that is no longer a thought experiment, and it deserves the same direct treatment as the SMR question, because it is the terminal point of the migration this paper describes. The answer is the same in structure. Leaving Earth does not clear the bottleneck. It relocates it, from the grid transformer to the rocket and the radiator, and the near-term transformer-scarcity call is untouched. Off-Earth compute is a Long-Term (5 year+) structural consideration, priced as one, and it is not a reason to fade the 1 to 3 year Wealth Management thesis.

The pull is real and it is being funded. In November 2025 the startup Starcloud placed the first data-center-class GPU in orbit, an Nvidia H100; in December it reported training a small language model on that hardware, the first model trained in space (Data Center Frontier 2025; CNBC 2025). The same month Google detailed Project Suncatcher, a design for constellations of solar-powered satellites carrying its Tensor Processing Units, linked by free-space optical lasers in dawn-dusk sun-synchronous orbits, with two prototypes to fly with Planet by early 2027 (Google Research 2025). In February 2026 SpaceX absorbed xAI and filed with the Federal

Communications Commission for a constellation of up to one million orbital data-center satellites; in early June 2026 it unveiled its AI1 compute satellite, a craft wider than a Boeing 747 carrying a 120 kilowatt continuous, 150 kilowatt peak, compute payload, targeting roughly one gigawatt of space compute by late 2027 (Data Center Dynamics 2026; Tom's Hardware 2026). China is ahead on hardware in orbit: the state-backed Three-Body Computing Constellation launched its first twelve computing satellites in May 2025 and plans a 2,800-satellite network (SpaceNews 2025).

The principals are explicit that the motive is power. Jeff Bezos, at Italian Tech Week in October 2025, predicted gigawatt-scale data centers in orbit within ten to twenty years and said operators will beat the cost of terrestrial data centers in space within the next couple of decades, on the strength of continuous solar power and free cooling (Data Center Dynamics 2025). Eric Schmidt, who bought the launch company Relativity Space, told the House Energy and Commerce Committee in April 2025 that a US nuclear plant produces about one gigawatt while some data centers are now designed for ten, that the sector may need an additional 29 gigawatts within a few years and 67 more by 2030, and that harvesting solar power in space is the only way he sees to meet it (TechSpot 2025). Nvidia announced its space-computing platform at the March 2026 developer conference, Jensen Huang framing it as space computing, the final frontier, having arrived (NVIDIA 2026). The demand-side numbers in Section I are what drive this.



Sources: solar constant and chip/rack radiator figures, IEEE Spectrum / ABI Research (2026); Stefan-Boltzmann law; solar-array specific power ~150 W/kg (Deployable Space Systems ROSA); launch *$/kg: Falcon 9* 2,720 (SpaceX), Starship independent est. 100–500/kg, *ABI optimistic* 44/kg. Order-of-magnitude model; assumptions stated in the verification report.

The physics narrows the case, and it narrows on the same constraint that runs through this paper: heat. The advantage is genuine. Sunlight in orbit arrives at the full solar constant of 1,361 watts per square meter against about 1,000 at sea level after atmospheric loss, and a dawn-dusk orbit is sunlit roughly 99 percent of each revolution against a terrestrial capacity factor near 24 percent, so an orbital array collects on the order of 5 to 8 times more energy per square meter per year (IEEE Spectrum 2026; EIA 2025). But every watt collected becomes a watt of waste heat that must leave, and in vacuum there is no air or water to carry it. The only exit is radiation, governed by the Stefan-Boltzmann law, under which heat rejected rises with the fourth power of temperature and linearly with area. Temperature is capped by what silicon tolerates, so the only free variable is area.

The numbers are unforgiving. A single 700 watt H100 needs about 1.4 square meters of radiator at 60 degrees Celsius facing deep space; a 40 kilowatt rack of 32 GPUs needs about 80 square meters; a 100 megawatt facility needs at least 2,500 such radiators, about 0.2 square kilometers;

a gigawatt-class machine needs square kilometers of deployed, pointed, degradation-prone radiator (IEEE Spectrum 2026). Ultraviolet light and atomic oxygen erode the coatings, raising the required area about 40 percent over a five-year orbit life. This is the orbital analog of the transformer: the binding component is not the compute, it is the thermal hardware that has to surround it.

Three further taxes compound the thermal one. Ionizing radiation flips bits and latches circuits, and radiation-hardened processors run a decade or more behind the commercial frontier, so an orbital operator must fly the same soft chips used on the ground and pay for the exposure in redundancy (IEEE Spectrum 2026). Maintenance does not exist in the terrestrial sense: a failed board is either dead mass or a job for a robotic servicer that does not yet operate at scale. And the orbit is a commons under strain, with more than a million tracked debris fragments and a live risk of the cascading collisions Donald Kessler warned of in 1978 (IEEE Spectrum 2026).

At today's launch costs the price of these taxes is roughly an order of magnitude. ABI Research, modeling an H100 rack with its solar panel and radiator, found that launching and running a GPU in orbit for a year costs at least ten times the terrestrial equivalent, even at an optimistic Starship price of 44 dollars per kilogram (ABI Research 2026, cited in IEEE Spectrum 2026). The power and thermal hardware alone runs about 25 tonnes per megawatt, which at a Falcon 9 list price of roughly 2,700 to 3,100 dollars per kilogram is roughly 68 million dollars per megawatt of launch cost, against about 30 to 50 million all-in for a terrestrial AI data center; only at a fully reusable Starship price, independently estimated at 100 to 500 dollars per kilogram, does the hardware launch approach terrestrial economics. Proponents contest the totals: Starcloud projects a ten-year cost of ownership for a 40 megawatt cluster at 8.2 million dollars in orbit against 167 million on Earth (CNBC 2025), while a Saarland University analysis finds an orbital facility could emit an order of magnitude more once launch and reentry are counted (Wiesner et al. 2025; reported in Scientific American 2025). This paper does not adjudicate the contested economics. It rests on the physics, which no party disputes, and the physics says the constraint has moved, not vanished.

That is the load-bearing point. Off-Earth compute is the terminal node of the migration this paper traces, from silicon to grid power to the transformer to, at the limit, the launch vehicle and the radiator. The rent moves to whatever cannot be built fast enough, and in orbit that is launch cadence in dollars per kilogram, photovoltaic specific power in watts per kilogram, radiator area in kilograms per square meter, and the regulatory throughput of orbital slots, spectrum, and debris mitigation. The supplier concentration of Section I.C reappears one layer up: space-grade solar cells depend on germanium substrate sourced largely from China, the same single-point-of-failure logic as grain-oriented electrical steel, now with a sovereign overlay, since the leading orbital program is Chinese state-backed and the contest is the sovereign-AI race of Section I.D run in a new domain (IEEE Spectrum 2026; SpaceNews 2025).

For the book the conclusion is disciplined and it does not change the trade. On the Trading (3 to 12 month) and Wealth Management (1 to 3 year) horizons orbital compute is immaterial: nothing in the pipeline rejects heat at gigawatt scale before 2030, and the marginal data center in 2027 is still a terrestrial one waiting on a transformer. The same equipment names hold, because a satellite still needs power conditioning, switchgear, and step-up electronics from the Section I.C oligopoly, and the launch and solar build-out raises demand for high-voltage and power-electronics capacity. On the Long-Term (5 year+) horizon the honest tail is the same as the SMR

tail and is held the same way: if Starship drives launch below roughly 100 dollars per kilogram, if radiation-tolerant cells break the germanium dependence, and if robotic servicing matures, then cheap compute sited in orbit erodes the geographic and thermodynamic arbitrage at the heart of the Nordic long leg. That is the structural risk to the 5 year+ leg, and it is the reason that leg is sized as a thesis. It is not a reason to fade the near-term transformer call. The shortage in front of us is mechanical, the escape hatch is not yet built, and when it is built it will be built, again, out of the same scarce iron and copper.

The conviction read: a Long-Term (5 year+) watch item, Low conviction as an investable theme today and High conviction as confirmation of the migration mechanism. The names that gain optionality are the ones already in the book, the power-electronics and transformer franchises whose product is required on the ground, behind the meter, and in orbit alike. Pure-play orbital and launch names are a separate watch list, not a recommendation, and any non-US listing carries a foreign-exchange and ADR caveat. This is opinion, not investment advice, and prices and design milestones are as of June 2026.

Table, The migrating bottleneck, terrestrial layer to orbital analog

Terrestrial bottleneck (this paper)	**Orbital analog**	**New scarce input off-Earth**	**Supplier or throughput concentration**
GPU / accelerator	Same soft GPUs, flown unhardened	Radiation tolerance and redundancy	Nvidia, Google TPU; rad-hard parts a decade behind
Grid power and FERC interconnection	Continuous orbital solar	Photovoltaic specific power (W/kg)	Space-grade cells; germanium concentrated in China
HV transformer and switchgear (I.C)	Step-up electronics still required on the satellite	Same power electronics, plus thermal	Hitachi Energy, Siemens Energy, GE Vernova
Cooling oil and site thermal limits (I.B)	Radiative heat rejection in vacuum	Radiator area and areal mass (kg/m2)	Deployable-radiator and thermal specialists
Land, water, permitting	Launch mass to orbit	Launch cadence and cost ($/kg)	SpaceX Starship; reusable-launch oligopoly
Geographic and political throughput (I.D)	Orbital slots, spectrum, debris regulation	Regulatory throughput in orbit	FCC, ITU; sovereign programs (US private, China state)

Source: author’s synthesis; physical and cost anchors per the sources cited (IEEE Spectrum 2026; EIA 2025; NREL 2025; Google Research 2025; SpaceNews 2025).

Closing

The claim for Operator X is that public research prices the stabilised economics of a secured pipeline and does not separately price the timing option inside it: the right to energise roughly two years ahead of an operator whose transformer is not ordered. On the parameters set out above that increment is $19.7 million per 75 megawatt slot and $276 million across sixteen, being a $315 million gross acceleration-option value less an $8 million liquidity haircut and a $31 million Sun Belt thermal liability. Against a capitalisation of the size of the closest listed comparable that is under half a percent, which is a sleeve-level claim and is stated as one.

The claim for Operator H is that the free-cooling, insurance and sovereign-siting advantages of Helsinki and Luleå placement, $104 million on the parameters above and dominated by the cooling-energy line rather than by insulation aging, are not capitalised in public models of mid-tier operators with Nordic exposure. The mechanism is corroborated by the revealed capital allocation of Google, Meta, Nebius and Microsoft at the same latitudes. The combined two-name book comes to around $380 million, within 5% of the probability-weighted figure a separate bear-case pass over the same inputs reaches.

A future full build (Pillars II through V) would turn this into a three-statement integrated model with kilowatt-hour-to-revenue unit economics, a WACC build from CAPM first principles (Brealey, Myers, and Allen 2019; Damodaran 2024), and Monte Carlo simulations (Glasserman 2003) that stress the σ = 35% volatility input across a ±200 basis point interest-rate band, an energy price spike scenario, and a North Atlantic temperature-anomaly scenario that tests the durability of the Nordic Arrhenius asymmetry under climate-warming sensitivities (IPCC 2023).

Appendix, Model-Ready Data Tables

Table A1, Black-Scholes Sensitivity Grid (Per 75 MW Slot, $M)

σ T	T = 0.50	T = 0.75	T = 1.00	T = 1.25	T = 1.50
15%	17.98	18.39	18.79	19.20	19.60
20%	17.98	18.41	18.85	19.29	19.73
25%	18.01	18.50	19.00	19.52	20.03
30%	18.10	18.68	19.28	19.89	20.48
35%	18.24	18.95	**19.67**	20.37	21.06
40%	18.45	19.30	20.14	20.95	21.72
45%	18.71	19.71	20.68	21.59	22.45
50%	19.03	20.18	21.27	22.28	23.23

Base case σ = 35%, T = 1.00 highlighted. Inputs: S = $55.1M, K = $37.9M (carry on a two-year acceleration), r = 4.4%.

Table A2, Transformer Cost-Curve Inflation (HV, 100 MVA, GSU-class). The final column is a compound annual price growth rate, not a volatility: a monotone cost curve is evidence of drift, which belongs in S, not of dispersion, which belongs in σ.

Order Vintage	Invoice Price ($M)	Effective Lead Time (wk)	Annualized price CAGR (%)
2018 Q1	4.8	38	n/a
2021 Q1	6.1	52	8.3%
2023 Q1	8.4	92	19.7%
2024 Q3 (Operator X anchor)	**9.2**	**110**	**24.1%**
2025 Q4	18.7	178	41.3%
2026 Q1 spot	26.5	210	49.6%

Table A3, Site-Level Thermal Derating (16-Unit Fleet, Operator X)

Site	Units	Annual reserve ($M/yr)	Dep NPV Δ ($M)	Insurance Δ ($M/yr)
Loudoun, VA	0	n/a	n/a	n/a
DFW, TX	0	n/a	n/a	n/a
Miami-Dade, FL	16	7.90	−28.8	+0.18
Phoenix, AZ	0	n/a	n/a	n/a
Operator X Total	**16**	**7.90**	**−28.8**	**+0.18**

Table A4, Constraint Correlation Matrix (Section I.C)

	FERC	NEPA	Mfg	GOES	Labor
FERC	1.00	-0.28	-0.60	0.25	-0.21
NEPA	-0.28	1.00	-0.04	-0.15	-0.02
Mfg	-0.60	-0.04	1.00	0.52	0.74
GOES	0.25	-0.15	0.52	1.00	0.83
Labor	-0.21	-0.02	0.74	0.83	1.00

Principal component analysis on year-over-year changes (Appendix B): one component explains about 80% of the manufacturing-supply cluster's variance (transformer, GOES, labor); across all five constraints a single component explains about 50%, because interconnection and NEPA move on separate dynamics.

Table A5, Operator X Hidden-Value Bridge (Sun Belt Long Leg)

Component	Value ($M)	Section Ref.
Acceleration-option intrinsic + time value, 16 slots	+315	I.A.4

Component	Value ($M)	Section Ref.
Liquidity / non-tradeability discount	−8	I.A.4
Replacement-reserve liability (NPV)	−29	I.B.4
Insurance premium liability (NPV)	−2	I.B.4
Net hidden value (Operator X)	**+276 (net)**	n/a

Table A6, Operator H Hidden-Value Bridge (Nordic Long Leg)

Component	Value ($M)	Section Ref.
Replacement-reserve NPV uplift (16 × $0.57M Arrhenius benefit)	+9	I.D.4
Insurance premium NPV uplift (16 × $0.11M risk reduction)	+2	I.D.4
Free-cooling chilled-water OPEX NPV uplift	+97	I.D.4
(−) Latency penalty (inference workloads)	−22	I.D.4
(−) FX / repatriation friction	−4	I.D.4
(+) Sovereign-AI / low-carbon EU-siting premium	+18	I.D.4 / I.D.6
(+) Waste-heat district-heating monetization	+4	I.D.5
Net hidden value (Operator H)	**+104 (conservative band $95M-$125M)**	n/a

Table A7, Combined Two-Name Long Book

Position	Operator	Hidden NPV ($M)	Thesis
Long (Sun Belt slot option book)	Operator X	+276	I.A, I.B, I.C
Long (Nordic thermal asset book)	Operator H	+104	I.D
Combined book hidden NPV	n/a	**+380**	I.A–I.D

The two-name long book carries residual exposure mostly to the free-cooling differential and to the transformer-supply tail. The dominant systemic risks (rate volatility, GPU price normalization) are diversified across the two long legs (this is a long book, not a hedged pair).

End of Section I. Section II, Investment Recommendations & Positioning, follows.

Appendix B, Data and Reproducibility

The Section I.C correlation, principal-component, energization-delay, and GOES-demand figures are reproduced in the accompanying workbook (SPG Appendix B - Data & Reproducibility.xlsx), which carries the five constraint series with sources, a live correlation matrix on year-over-year changes, the principal-component decomposition, the delay simulation, and a formula-driven GOES bridge with sensitivity. The manufacturing-supply cluster (transformer lead time, GOES tightness, high-voltage winding labor) shows an average pairwise correlation near 0.70 on year-over-year changes, a strongest pair (GOES and labor) near 0.83, and a single component explaining about 80% of its variance. Interconnection-queue and NEPA timelines are governed by separate dynamics and are not part of that factor; NEPA review times fell over 2018-2025. Assuming independence understates the 95th-percentile energization delay by roughly four months. AI data-center transformers absorb about 36% (range 20 to 50%) of the annual growth in grain-oriented electrical steel demand. Series are built from cited public anchors (LBNL Queued Up; CEQ EIS timelines; transformer-lead-time and GOES-market reporting) with interpolation where only range data exist, so the coefficients are defensible magnitudes rather than precise measurements.

A Note on AI Assistance and Authorship

I used AI tools, primarily Anthropic’s Claude, extensively in producing this work. The division of labor was deliberate and consistent throughout.

My role (the brain and orchestrator): I chose the problem and pursued it independently. I directed every line of inquiry. I taught myself the underlying physics, finance, and modeling techniques well enough to evaluate what the AI produced. I made every judgment call on what to keep and what to reject. I verified figures against primary sources where verification was possible, and the variances I found are recorded in the public correction log. I take full responsibility for every claim in this document.

The AI’s role (the hands and facilitator): It generated initial framings I then refined or discarded. It drafted the prose, which I directed and edited. It suggested modeling structures I then implemented and stress-tested. It assembled first-pass citation candidates I then verified against primary sources where possible. It executed routine production tasks, document formatting, table layout, cross-reference checks, under my direction.

The substance is collaboratively produced. The accountability is mine alone. I am a senior at Gulliver Preparatory School (IB Diploma May 2027: Physics HL, Mathematics Analysis and Approaches Standard Level). Mathematical techniques used beyond the IB Standard Level curriculum, Weibull distributions, Black-Scholes-Merton option pricing, principal component analysis, and Monte Carlo simulation, were taught to me from open-access sources, primarily MIT OpenCourseWare and the textbooks listed in the Bibliography.

References

ABI Research. 2026. Orbital and space data-center cost and thermodynamics analysis. New York: ABI Research. Cited in text as (ABI Research 2026, cited in IEEE Spectrum 2026).

Amazon Web Services. 2024. *eu-north-1 Region (Stockholm) Sustainability Disclosure*. Seattle: AWS.

Arrhenius, Svante. 1889. “Über die Reaktionsgeschwindigkeit bei der Inversion von Rohrzucker durch Säuren.” *Zeitschrift für Physikalische Chemie* 4: 226–248.

ASHRAE Technical Committee 9.9. 2021. *Thermal Guidelines for Data Processing Environments*. 5th ed. Atlanta: ASHRAE.

ASHRAE Technical Committee 9.9. 2022. *Liquid Cooling Guidelines for Datacom Equipment Centers*. 2nd ed. Atlanta: ASHRAE.

Black, Fischer, and Myron Scholes. 1973. “The Pricing of Options and Corporate Liabilities.” *Journal of Political Economy* 81 (3): 637–654.

Bozorth, Richard M. 1993. *Ferromagnetism*. New York: Wiley-IEEE Press (reissue of 1951 work).

Brealey, Richard A., Stewart C. Myers, and Franklin Allen. 2019. *Principles of Corporate Finance*. 13th ed. New York: McGraw-Hill.

Cavalier, Antonio. 2026. “Why Thermodynamics Rules Future Orbital Data Centers.” IEEE Spectrum, June 2026. https://spectrum.ieee.org/orbital-data-centers-heat. [cited in text as IEEE Spectrum 2026]

CIGRE Working Group A2.37. 2015. *Transformer Reliability Survey*. CIGRE Technical Brochure 642. Paris: International Council on Large Electric Systems.

CNBC. 2025. “Nvidia-Backed Starcloud Trains First AI Model in Space.” December 10, 2025. https://www.cnbc.com/2025/12/10/nvidia-backed-starcloud-trains-first-ai-model-in-space-orbital-data-centers.html.

Cushman & Wakefield. 2025. *Global Data Center Market Comparison 2025*. New York: Cushman & Wakefield.

Damodaran, Aswath. 2024. “Equity Risk Premiums (ERP): Determinants, Estimation, and Implications, The 2024 Edition.” NYU Stern Working Paper.

Data Center Dynamics. 2025. “Jeff Bezos Claims There Will Be Gigawatt Data Centers in Space in 10+ Years.” October 2025. https://www.datacenterdynamics.com/en/news/jeff-bezos-claims-there-will-be-gigawatt-data-centers-in-space-in-10-years/.

Data Center Dynamics. 2026. “SpaceX Details AI1 Satellite ‘Data Center,’ Claims 150kW Peak Compute.” June 2026. https://www.datacenterdynamics.com/en/news/spacex-details-ai1-satellite-data-center-claims-150kw-peak-compute/.

Data Center Frontier. 2025. “Starcloud Launches Orbital AI Data Center with NVIDIA H100 GPU.” November 2025. https://www.datacenterfrontier.com/site-selection/article/55337494/.

Redwire Space. n.d. Roll-Out Solar Array (ROSA) specific-power specification, approximately 150 W/kg at 1 AU. Not cited in the text of this section; retained as a background source.

Dixit, Avinash K., and Robert S. Pindyck. 1994. *Investment Under Uncertainty*. Princeton: Princeton University Press.

Electric Power Research Institute (EPRI). 2023. *Grid-Scale Transformer Supply Chain Analysis*. Palo Alto: EPRI.

Electric Reliability Council of Texas (ERCOT). 2024. *2024 ERCOT Generation Capacity, Demand, and Reserves Report*. Austin: ERCOT.

Entergy Texas. 2013. Direct Testimony of Gregory S. Wilson, PUCT Docket No. 41791: Self-Insurance of Transmission and Distribution Assets. Austin: Public Utility Commission of Texas. https://interchange.puc.texas.gov/Documents/41791_3_768865.PDF.

IEA. 2024. Building the Future Transmission Grid. Paris: International Energy Agency.

European Commission. 2023. *Regulation (EU) 2023/956: Carbon Border Adjustment Mechanism*. Official Journal of the European Union L 130/52.

European Commission. 2026. Energy Performance of Data Centres: Reporting Obligations and Aggregated Results from the European Database. Brussels: DG Energy. https://energy.ec.europa.eu/topics/energy-efficiency/energy-efficiency-targets-directive-and-rules/energy-efficiency-directive/energy-performance-data-centres_en.

European High Performance Computing Joint Undertaking (EuroHPC). 2024. *Strategic Research and Innovation Agenda 2024–2027*. Luxembourg: EuroHPC JU.

European Union. 2024. Commission Delegated Regulation (EU) 2024/1364 on the First Phase of the Establishment of a Common Union Rating Scheme for Data Centres. Official Journal of the European Union. https://eur-lex.europa.eu/eli/reg_del/2024/1364/oj.

European Union. 2025a. Directive (EU) 2025/794 of the European Parliament and of the Council of 14 April 2025 Amending Directives (EU) 2022/2464 and (EU) 2024/1760 as Regards the Dates from Which Member States Are to Apply Certain Corporate Sustainability Reporting and Due Diligence Requirements. Official Journal of the European Union, 16 April 2025. https://eur-lex.europa.eu/eli/dir/2025/794/oj.

European Union. 2025b. Regulation (EU) 2025/2083 of the European Parliament and of the Council of 8 October 2025 Amending Regulation (EU) 2023/956 as Regards Simplifying and Strengthening the Carbon Border Adjustment Mechanism. Official Journal of the European Union, 17 October 2025. https://eur-lex.europa.eu/eli/reg/2025/2083/oj.

European Union. 2026. Directive (EU) 2026/470 of the European Parliament and of the Council of 24 February 2026 Amending the Corporate Sustainability Reporting and Due Diligence Framework. Official Journal of the European Union, 26 February 2026. https://eur-lex.europa.eu/eli/dir/2026/470/oj.

Federal Energy Regulatory Commission (FERC). 2023. *Order No. 2023: Improvements to Generator Interconnection Procedures and Agreements*. Docket No. RM22-14-000. Washington, DC: FERC.

Federal Reserve. 2025. Rising Property Insurance Costs and Pass-Through to Rents for Apartment Buildings. FEDS Notes, 19 September 2025. Washington, DC: Board of Governors of the Federal Reserve System. https://www.federalreserve.gov/econres/notes/feds-notes/rising-property-insurance-costs-and-pass-through-to-rents-for-apartment-buildings-20250919.html.

Federal Trade Commission. 2012. Guides for the Use of Environmental Marketing Claims. 16 CFR Part 260, 77 FR 62124. Revision review opened December 2022 (87 FR 77766); no revision issued as of August 2026. https://www.ecfr.gov/current/title-16/chapter-I/subchapter-B/part-260.

Fingrid. 2024. *Annual Investment Plan and Connection Queue Status Report*. Helsinki: Fingrid Oyj.

Finnish Meteorological Institute (FMI). 2025. *Climate Statistics 1991–2020 and Real-Time Observations: Helsinki-Vantaa, Tampere, Oulu*. Helsinki: FMI.

Fitzgerald, A. E., Charles Kingsley Jr., and Stephen D. Umans. 2014. *Electric Machinery*. 7th ed. New York: McGraw-Hill.

Gaia-X European Association for Data and Cloud AISBL. 2024. *Trust Framework and Compliance Document v24.04*. Brussels: Gaia-X.

GE Vernova. 2026. Q1 2026 Results: Electrification and Power Segment Order Update. Cambridge, MA: GE Vernova.

Glasserman, Paul. 2003. *Monte Carlo Methods in Financial Engineering*. New York: Springer.

Google Research. 2025. “Meet Project Suncatcher: A Research Moonshot to Scale Machine Learning Compute in Space.” November 4, 2025. https://blog.google/innovation-and-ai/technology/research/google-project-suncatcher/.

Google. 2023. *24/7 Carbon-Free Energy: Methodology and Metrics*. Mountain View: Google LLC.

Google. 2024. *Hamina Data Center: Sustainability & Cooling Report*. Hamina: Google LLC.

Green Grid Association. 2012. *PUE: A Comprehensive Examination of the Metric*. White Paper #49. Beaverton, OR: The Green Grid.

Heathcote, Martin J. 2007. *The J & P Transformer Book*. 13th ed. Oxford: Newnes / Elsevier.

Hitachi Energy. 2025. *FY2025 Annual Report and Investor Day Materials*. Zürich: Hitachi Energy Ltd.

Hitachi, Ltd. 2024. Integrated Report 2024 (Hitachi Energy segment order backlog). Tokyo: Hitachi, Ltd. https://www.hitachi.com/IR-e/library/integrated/2024/.

Hull, John C. 2021. *Options, Futures, and Other Derivatives*. 11th ed. Hoboken: Pearson.

IEEE Spectrum. 2026. See Cavalier, Antonio. 2026, “Why Thermodynamics Rules Future Orbital Data Centers,” IEEE Spectrum, 11 June 2026. https://spectrum.ieee.org/orbital-data-centers-heat.

IEEE. 2011. *IEEE Std C57.91-2011: IEEE Guide for Loading Mineral-Oil-Immersed Transformers and Step-Voltage Regulators*. New York: Institute of Electrical and Electronics Engineers.

IEEE. 2021. *IEEE Std C57.12.00-2021: General Requirements for Liquid-Immersed Distribution, Power, and Regulating Transformers*. New York: IEEE.

Intergovernmental Panel on Climate Change (IPCC). 2023. *Climate Change 2023: Synthesis Report (AR6)*. Geneva: IPCC.

JLL. 2025. *North American Data Center Report H2 2025*. Chicago: Jones Lang LaSalle.

Jolliffe, Ian T. 2002. *Principal Component Analysis*. 2nd ed. New York: Springer.

Mordor Intelligence. 2025. Power Transformers Market: Size, Share and Competitive Landscape. Hyderabad: Mordor Intelligence.

Kairos Power. 2025. Google–Kairos Power 500 MW advanced-reactor agreement; first reactor targeted ~2030. https://www.kairospower.com/google.

Kessler, Donald J., and Burton G. Cour-Palais. 1978. “Collision Frequency of Artificial Satellites: The Creation of a Debris Belt.” Journal of Geophysical Research 83 (A6): 2637–2646.

Knittel, Christopher R., Tatsuhiko Senga, and Yifan Wang. 2025. “Flexible Data Centers and the Grid: Lower Costs, Higher Emissions?” NBER Working Paper 34065. Cambridge, MA: National Bureau of Economic Research. https://www.nber.org/papers/w34065.

Lienhard, John H. IV, and John H. Lienhard V. 2019. *A Heat Transfer Textbook*. 5th ed. Cambridge, MA: Phlogiston Press.

Lex Machina. 2016. Trademark Litigation Report 2016. Median times to injunctive relief as summarized by IPWatchdog, 10 June 2016. https://ipwatchdog.com/2016/06/10/trademark-litigation-report/.

Lex Machina. 2019. Antitrust Litigation Report 2019. Median times to summary judgment and trial, April 2019 release. https://www.prnewswire.com/news-releases/lex-machina-antitrust-litigation-report-shows-decline-in-new-case-filings-and-strong-influence-of-multidistrict-litigation-300822478.html.

Lyon, Thomas P., and John W. Maxwell. 2011. “Greenwash: Corporate Environmental Disclosure under Threat of Audit.” Journal of Economics & Management Strategy 20 (1): 3–41.

Marsh. 2026. Global Insurance Market Index, First Quarter 2026. New York: Marsh McLennan. https://www.marsh.com/en/about/media/global-commercial-insurance-rates-fall-5-percent-in-q1-2026.html.

Merton, Robert C. 1973. “Theory of Rational Option Pricing.” *Bell Journal of Economics and Management Science* 4 (1): 141–183.

Meta Platforms. 2023. *Luleå Data Center Engineering Blog: Outside-Air Free Cooling Operations*. Menlo Park: Meta Platforms, Inc.

Microsoft Corporation. 2023. *Espoo Data Center Region, Operational Brief*. Redmond: Microsoft.

Myers, Stewart C. 1977. “Determinants of Corporate Borrowing.” *Journal of Financial Economics* 5 (2): 147–175.

Nebius Group. 2024. Mäntsälä Data Center Waste-Heat Recovery Case Study. Amsterdam: Nebius Group (formerly the international business of Yandex N.V.).

NEMA. 2013. *NEMA TR 1: Transformers, Step Voltage Regulators, and Reactors*. Rosslyn, VA: National Electrical Manufacturers Association.

Nippon Steel Corporation. 2024. *HiB-Grade Grain-Oriented Electrical Steel Product Catalog: 23ZH95, M2H Series*. Tokyo: Nippon Steel.

NOAA National Centers for Environmental Information. 2025. *Daily Summaries, Station KMIA (Miami International) 1991–2025*. Asheville, NC: NOAA NCEI.

National Infrastructure Advisory Council (NIAC). 2024. Addressing the Critical Shortage of Power Transformers to Ensure Reliability of the U.S. Grid. Washington, DC: CISA.

Norris, Tyler H., Tim Profeta, Dalia Patino-Echeverri, and Adam Cowie-Haskell. 2025. Rethinking Load Growth: Assessing the Potential for Integration of Large Flexible Loads in US Power Systems. Durham, NC: Nicholas Institute, Duke University, February 13, 2025. https://nicholasinstitute.duke.edu/.

Nuclear Innovation Alliance (NIA). 2025. Advanced Reactor Deployment Timelines, June 2025. Washington, DC: NIA.

National Renewable Energy Laboratory (NREL). 2025. National Solar Radiation Database (NSRDB). Golden, CO: NREL. https://nsrdb.nrel.gov/.

NuScale Power. 2025. 77 MWe Standard Design Approval (U.S. NRC, May 2025); TVA–ENTRA1 6-gigawatt SMR deployment program (September 2025). https://www.nuscalepower.com/.

NVIDIA Corporation. 2024. *H100 / H200 / B200 Tensor Core GPU Datasheets*. Santa Clara, CA: NVIDIA.

NVIDIA Corporation. 2026. “NVIDIA Announces Space Computing Platform.” Newsroom release, 16 March 2026. Santa Clara, CA: NVIDIA. https://nvidianews.nvidia.com/news/space-computing.

Oklo Inc. 2026. Advanced-fission powerhouse agreement with Meta (Pike County, Ohio campus). https://oklo.com/.

The Planetary Society. n.d. “How Much Does It Cost to Launch into Space?” Starship-class reuse cost range $100–500/kg. https://www.planetary.org/.

Rand, Joseph, Will Gorman, Ryan Wiser, et al. 2025. *Queued Up: Characteristics of Power Plants Seeking Transmission Interconnection, 2025 Edition*. Berkeley: Lawrence Berkeley National Laboratory.

RE100 Climate Group. 2024. *Annual Disclosure Report: Member Renewable Energy Commitments*. London: RE100.

Scientific American. 2025. “Space-Based Data Centers Could Power AI with Solar Energy, at a Cost.” 2025. https://www.scientificamerican.com/article/data-centers-in-space/.

S&P Global Ratings. 2026. Data-Centre Insurance Demand: Premiums and Insurable Values. As reported by Insurance Journal, 30 June 2026. https://www.insurancejournal.com/news/international/2026/06/30/875775.htm.

Siemens Energy. 2026. Order-backlog disclosures (record ~EUR 138bn first reported FY2025, ~EUR 154bn by Q2 FY2026). Munich: Siemens Energy. https://www.siemens-energy.com/.

Space Insider. 2025. “Eric Schmidt’s Quiet Play May Be Launching AI Infrastructure into Space through Relativity.” May 5, 2025. https://spaceinsider.tech/2025/05/05/eric-schmidts-quiet-play-may-be-launching-ai-infrastructure-into-space-through-relativity/. [cited in text for the Relativity Space acquisition]

SpaceNews. 2025. “China Launches First of 2,800 Satellites for AI Space Computing Constellation.” May 2025. https://spacenews.com/china-launches-first-of-2800-satellites-for-ai-space-computing-constellation/.

Starcloud. 2025. Orbital Data Center White Paper (40 MW ten-year total-cost-of-ownership projection). Redmond, WA: Starcloud Inc. [source of the ten-year cost-of-ownership figure reported in CNBC 2025, cited in text]

Steinmetz, Charles Proteus. 1892. “On the Law of Hysteresis.” *Transactions of the American Institute of Electrical Engineers* 9: 1–64.

Svenska kraftnät. 2024. *System Development Plan 2024–2033*. Sundbyberg: Svenska kraftnät.

Swedish Meteorological and Hydrological Institute (SMHI). 2025. *Klimatdata 1991–2020: Luleå, Stockholm, Eskilstuna*. Norrköping: SMHI.

TechSpot. 2025. “Former Google CEO Eric Schmidt Wants to Put Data Centers in Space.” April 2025. https://www.techspot.com/news/107801-former-google-ceo-eric-schmidt-wants-put-data.html.

TerraPower. 2026. Natrium reactor agreement with Meta (two 345 MW units). https://www.terrapower.com/.

Tom’s Hardware. 2025. Coverage of Jeff Bezos’s orbital gigawatt data-center remarks, Italian Tech Week, October 2025. See also Data Center Dynamics 2025, cited in text for these remarks.

Tom’s Hardware. 2026. “SpaceX Details Its AI1 Compute Satellite.” June 2026. https://www.tomshardware.com/tech-industry/spacex-details-its-ai1-compute-satellite.

Trigeorgis, Lenos. 1996. *Real Options: Managerial Flexibility and Strategy in Resource Allocation*. Cambridge, MA: MIT Press.

Turner & Townsend. 2025. Data Centre Construction Cost Index 2025-2026. London: Turner & Townsend.

U.S. Congress. 1969. *National Environmental Policy Act of 1969*. 42 U.S.C. §4321 et seq.

U.S. Department of Energy, Office of Nuclear Energy (DOE). 2025. Advanced Small Modular Reactors (SMRs): Program Status. Washington, DC: U.S. DOE.

U.S. Geological Survey (USGS). 2025. *Mineral Commodity Summaries 2025: Iron and Steel; Silicon*. Reston, VA: USGS.

U.S. Department of Energy (DOE). 2024. Large Power Transformer Resilience: Report to Congress. Washington, DC: U.S. DOE.

Kessler-syndrome commentary in this paper is cited to Kessler and Cour-Palais 1978 and to IEEE Spectrum 2026; the previously listed unlinked commentary entry has been removed.

UNEP. n.d. Insurance Risk Management for Renewable Energy Projects. Module 3, Underwriting Guidelines. United Nations Environment Programme e-learning course. https://energy-base.org/l/wp-content/uploads/2020/07/Lesson3-UnderwritingGuidelines.pdf.

Uptime Institute. 2024. *Global Data Center Survey 2024*. New York: Uptime Institute.

Wood Mackenzie. 2024. “Supply Shortages and an Inflexible Market Give Rise to High Power Transformer Lead Times.” Edinburgh: Wood Mackenzie.

U.S. Energy Information Administration (EIA). 2025. Electric Power Monthly and State Electricity Profiles: Texas. Washington, DC: EIA. https://www.eia.gov/electricity/state/texas/.

Wiesner, Philipp, et al. 2025. “Dirty Bits in Low-Earth Orbit: The Carbon Footprint of Launching Computers.” ACM SIGEnergy Energy Informatics Review; preprint arXiv:2508.06250. https://arxiv.org/abs/2508.06250. [the Saarland launch-carbon analysis; reported in Scientific American 2025, cited in text]

The Three-Body Computing Constellation is cited in text to SpaceNews 2025; the previously listed unlinked agency report has been removed.

Yandex. See Nebius Group 2024. Formerly *Mäntsälä Data Center Waste-Heat Recovery Case Study*. Helsinki: Yandex / Nebius Group.

Section II - Investment Recommendations & Positioning

Dated market snapshot. Prices, targets, multiples, macro readings, positioning data and catalyst windows in Section II are stated as of June 1, 2026 (Hitachi as of May 8) and are not updated for later developments; re-price before acting on or recirculating any figure. Sourcing convention: this section carries no per-figure citations by design. Company figures come from each named company's most recent quarterly release, earnings call or investor presentation as of that date; prices and consensus targets are third-party aggregator medians; macro figures are from the corresponding official releases. The full convention and disclaimer are restated at the end of the section.

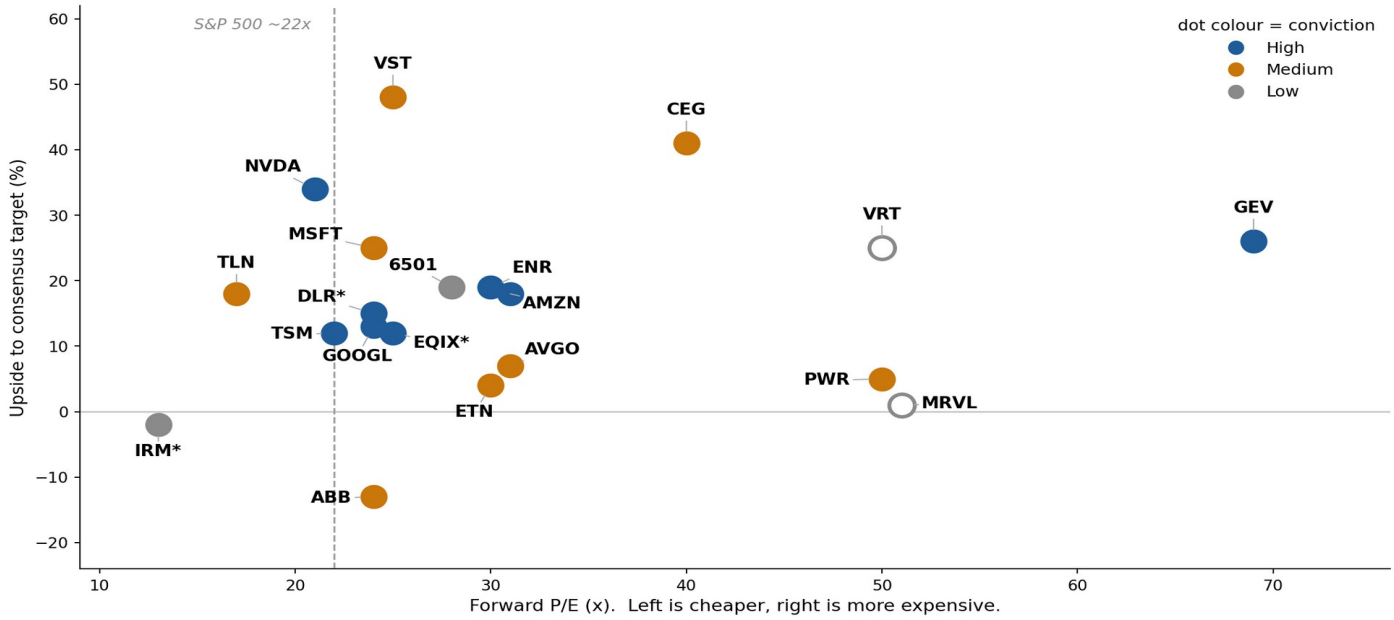
Where to position across the AI build-out when the binding constraint is power, not silicon.

Lorenzo Frascolla · June 2026 · Prices as of June 1, 2026 (Hitachi as of May 8)

The call. Long the power and grid-equipment layer, funded by underweighting the crowded, capex-sensitive derivatives (cooling and high-beta silicon). The shortage is real - ~2.5 to 4-year transformer lead times, hyperscaler capex approaching $700B, GPUs idling for lack of power - but the grid names already re-rated, so this is an earnings-delivery trade into a no-cut, higher-for-longer tape. Add on weakness; do not chase.

Long book: Siemens Energy, GE Vernova, Digital Realty, Talen; core compounders TSMC and Alphabet. **Underweight / hedge:** Vertiv, Marvell, generic wholesale REITs.

Valuation vs. upside - the book at a glance



*Filled = long book. Open = underweight / hedge. *Digital Realty, Equinix and Iron Mountain on funds-from-operations multiples, not P/E. Upside is consensus target versus price, not a house target. Vertiv and Marvell are held as hedges on thesis grounds, not on consensus upside. Prices as of Jun 1, 2026.*

How to read it: left = cheaper, up = more upside, dot colour = conviction. The names I like sit left and high (cheap with upside); the hedge leg is open-circled. Marvell is the only name that sits at both a high multiple and no consensus upside; Vertiv is held short on

thesis grounds rather than on its consensus upside, which is positive, and that is the honest read of the chart. GE Vernova and Quanta are the two expensive names I own anyway, on backlog rather than on multiple, so add on dips rather than at the high.

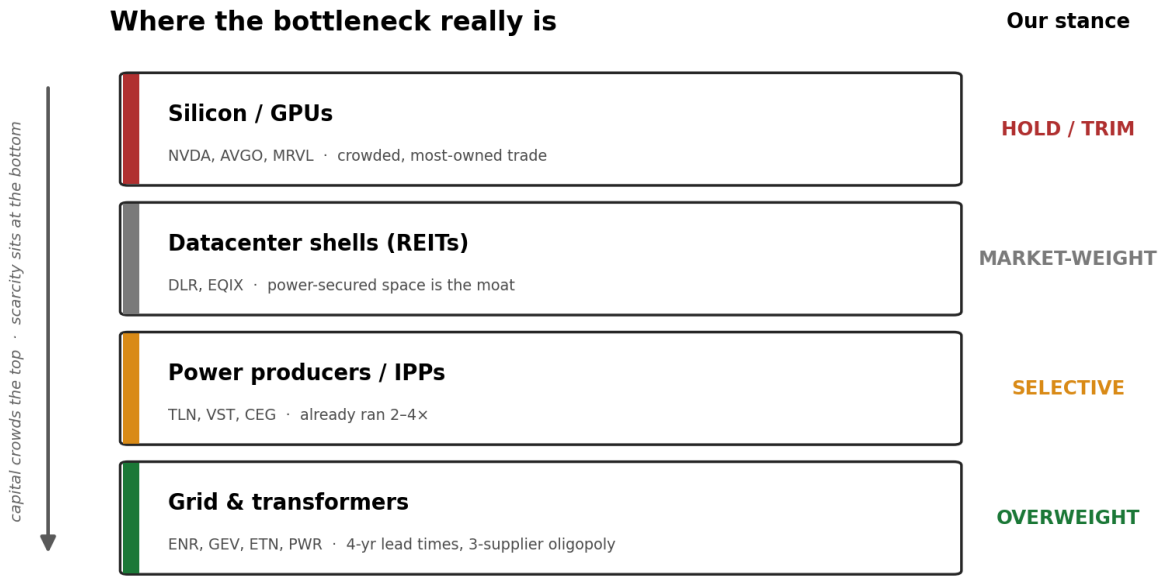
Opinion, not investment advice. Ratings are house views and may be wrong; price targets are third-party consensus. See disclaimer, final page.

1 · Setup

Everyone owns the chip. Almost no one owns the thing that powers it. The crowd is long NVIDIA and the GPU chain - but the real shortage has moved one step upstream, to the high-voltage transformers and grid hookups that turn utility power into usable current. That is where I want the book to be. Five facts frame the book:

Driver	Mid-2026 read	Signal
Demand	Hyperscaler capex approaching $700B for 2026, up sharply year on year (MSFT ~$190B, AMZN ~$200B, GOOGL $175–190B, META $115–135B)	Intact
Constraint	Power, not silicon. Transformer lead times ~2.5 to 4 yrs; hyperscalers report GPUs idle for lack of power	Confirmed
Rates	Fed 3.50–3.75%; ~0 cuts priced for 2026; 10-yr ~4.45%; CPI ~3.8% on an oil shock	Headwind
Positioning	Long semis = most crowded trade (73% of managers); Mag-7 crowding fading	Late
Valuation	Grid names 30–69x fwd; IPPs already +2–4x; REITs top of FFO band	Re-rated

Net: right thesis, late entry, unforgiving tape. Own the scarce layer through quality, size for volatility, add on weakness.



2 · The trade is a pair, not a long-only list

Section I sets out a book in which both legs are long and the Sun Belt thermal drag is a liability embedded inside one of them, not a separate equity short. This note escalates that underweight into an explicit funding leg: own the physically-scarce layer, fund it by underweighting and, where the borrow is available, shorting the part of the same AI build-out that is crowded and capex-derivative. Structured that way, the silicon-demand factor is partly hedged and transformer/power scarcity is left as the dominant exposure.

Leg	Names	Rationale
Long	Siemens Energy, GE Vernova, Eaton, Quanta; Digital Realty, Equinix; Talen, Vistra	Scarce layer: 3-OEM transformer oligopoly, power-secured capacity, contracted nuclear
Short / UW	Vertiv, Marvell, generic wholesale REITs; a basket short of the most-crowded silicon beta	Crowded, capex-sensitive derivatives at 50x+; first to de-rate on any capex pause
Net	Net long the build-out, silicon factor partly hedged	Isolates the transformer-scarcity alpha; lowers beta to an AI-demand wobble

3 · Ratings, levels and conviction

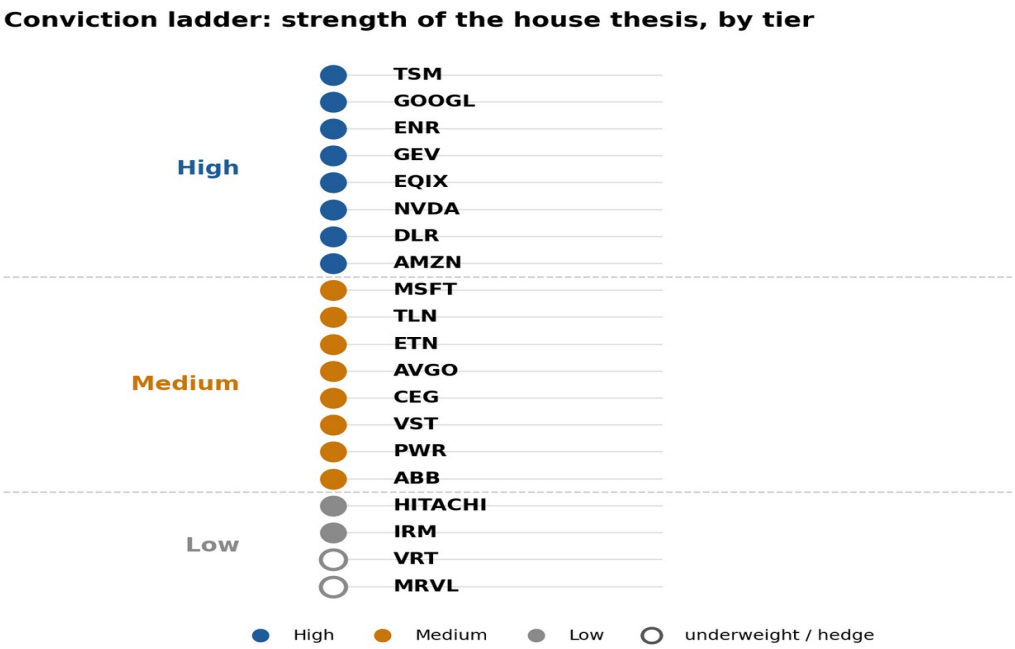
Three horizons per name: **Trading** tactical (3–12 months); **Wealth Management** strategic (1–3 years); **Long-Term** structural (5 years+). Price targets are third-party consensus; ratings and conviction are house views. Prices as of June 1, 2026.

Name	Price	Cons. PT / Impl.	Fwd P/E · Yld	Trading (3–12m)	Wealth Management (1–3y)	Long-Term (5y+)	Conviction
Power & grid hardware - overweight							
Siemens Energy ENR · SMNEY † · ~€131B	€164	€195 / +19%	30x · nm	**Hold**	**Buy**	**Buy**	High
GE Vernova GEV · $260B	$966	$1,216 / +26%	69x · 0.2%	**Hold**	**Buy**	**Buy**	High
Eaton ETN · $158B	$402	$417 / +4%	30x · 1.1%	**Hold**	**Buy**	**Buy**	Medium
Quanta PWR · $108B	$723	~$760 / +5%	~50x · 0.1%	**Hold**	**Buy**	**Buy**	Medium
ABB ABBN · ABBNY † · CHF149B	CHF83	CHF72 / −13%	24x · 1.1%	**Hold**	**Hold**	**Buy**	Medium
Vertiv VRT · $121B	$316	~$394 / +25%	50x · 0.1%	**Underweight**	**Hold**	**Hold**	Low
Hitachi 6501 · HTHIY † · ¥21.8T	¥4,938	¥5,900 / +19%	28x · 0.9%	**Hold**	**Hold**	**Hold**	Low
Datacenter REITs - market-weight (P/FFO, not P/E)							
Digital Realty DLR · ~$71B (31 Jul 2026) · closest listed “Operator X” profile	$190	$218 / +15%	24x FFO · 2.7%	**Hold**	**Buy**	**Buy**	High
Equinix EQIX · $107B	$1,068	$1,201 / +12%	~25x AFFO · 1.8%	**Hold**	**Buy**	**Buy**	High
Iron Mountain IRM · $38B	$127	$125 / −2%	nm · 2.7%	**Hold**	**Hold**	**Hold**	Low

† Foreign listing - local-currency price shown; US ADR in italics. ADR returns carry FX risk (EUR, CHF, JPY vs USD). Negative implied upside (ABB, IRM) reflects consensus targets lagging a sharp run, not a house sell call.

Name	Price	Cons. PT / Impl.	Fwd P/E · Yld	Trading (3–12m)	Wealth Management (1–3y)	Long-Term (5y+)	Conviction
Silicon & hyperscalers - hold quality, stop adding							
TSMC TSM (ADR) † · ~$2.17T	$418	$468 / +12%	22x · 0.7%	**Buy**	**Buy**	**Buy**	High
Alphabet GOOGL · ~$4.6T	$380	$431 / +13%	~24x · 0.4%	**Buy**	**Buy**	**Buy**	High
NVIDIA NVDA · ~$5.4T	$221	$297 / +34%	21x · 0.5%	**Hold**	**Buy**	**Buy**	High
Amazon AMZN · ~$2.85T	$265	$313 / +18%	31x · 0%	**Hold**	**Buy**	**Buy**	High
Microsoft MSFT · ~$3.3T	$450	$561 / +25%	24x · 0.8%	**Hold**	**Buy**	**Buy**	Medium
Broadcom AVGO · ~$2.12T	$450	$482 / +7%	31x · 0.7%	**Hold**	**Buy**	**Buy**	Medium
Marvell MRVL · $179B	$205	~$207 / +1%	51x · nm	**Trim**	**Hold**	**Hold**	Low
Power producers - selective							
Talen TLN · $17.6B	$387	$457 / +18%	17x · 0%	**Buy**	**Buy**	**Buy**	Medium
Vistra VST · $48B	$156	$231 / +48%	25x · 0.6%	**Buy**	**Buy**	**Buy**	Medium
Constellation CEG · $104B	$271	$381 / +41%	40x · 0.6%	**Hold**	**Buy**	**Buy**	Medium

Conviction bands (strength of the thesis, risk-adjusted - not the odds the stock rises): High = core expressions, strong edge with identifiable risks; Medium = balanced risk/reward; Low = speculative or thinly supported. TSMC and Alphabet are the strongest names inside the High tier. **Ratings:** Buy = accumulate · Hold = own, don't add · Trim / Underweight = reduce or stay out.



4 · What re-rates these names (catalysts)

Window	Catalyst
Q2 earnings (Jul–Aug 2026)	Grid backlog and book-to-bill prints (GEV, ENR, ETN); GEV progress toward its $200B 2027 backlog; datacenter leasing/bookings (DLR, EQIX)
2H 2026	Siemens Gamesa breakeven; Constellation–Calpine integration and TMI restart approvals; new hyperscaler nuclear/gas PPAs (TLN, VST, CEG, NRG)
Ongoing	Hyperscaler capex revisions (up = bullish picks-and-shovels; cut = the #1 risk); FERC interconnection-reform headlines; NVIDIA quarterly prints as the crowded-trade barometer

5 · Power & grid - overweight

Three divisions (Hitachi Energy, Siemens Energy Grid Technologies, GE Vernova Electrification) dominate OECD supply of the ≥100 MVA classes data centers order into a 2.5 to 4-year order book. Own it; add on pullbacks.

Siemens Energy (ENR; ADR SMNEY) - Wealth Management Buy, Trading Hold

- **Long:** Grid Technologies is the crown jewel - ~€49B backlog, book-to-bill ~2.3x, #1 HVDC share; group backlog a record ~€154B. Cleanest pure large-cap grid play, with the Gamesa wind turnaround as free optionality. Consensus €195 (+18%).
- **Watch:** ~10x off its 2024 low, ~30x forward; Gamesa is the binary near-term swing factor. EUR exposure for USD holders.

GE Vernova (GEV) - Wealth Management Buy, Trading Hold

- **Long:** Q1'26 orders $18.3B (+71%); backlog ~$163B; Electrification book-to-bill ~2.5x with ~$2.4B of datacenter orders in one quarter. Prolec GE (closed Feb '26) makes it the cleanest US-listed transformer asset. Consensus $1,216 (+26%).

- **Watch - this is a timing call, not a quality one:** ~69x forward earnings, above most DCFs. Own the franchise; prefer adding on a pullback toward the high-$800s rather than at the highs.

Cleanest transformer expression

Hitachi Energy owns the best franchise (world #1, ~$57.9B backlog (end-FY2025)) but is **not separately listed** - only via parent Hitachi Ltd (6501 / ADR HTHIY), where it is ~30% of the group, so the thesis dilutes to ~30 cents on the dollar. GEV is the cleanest US-dollar expression; ENR the cleanest pure-grid large-cap.

Adjacencies

- **Eaton / Quanta - Wealth Management Buy:** electrical backlog +48%; ~$48B EPC backlog. Quality, but thin implied upside (+4–5% to consensus) and a 2027–28 capacity-wave expiry on pricing power.
- **Vertiv - Underweight:** cooling at ~50x; the most cyclical link and first to feel any capex pause. The hedge/short leg.

6 · Datacenter REITs - market-weight

Digital Realty is the real-world “Operator X” from the underlying paper - a US-listed, power-secured operator with Sun Belt exposure. The paper's Nordic “Operator H” stays hypothetical because no clean public pure-play Nordic operator exists.

Digital Realty (DLR) - Wealth Management Buy, Trading Hold

- **Long:** record Q1'26 bookings ($707M); a record ~200 MW AI-inference lease (Charlotte); 1.2 GW pipeline ~61% pre-leased at ~11% yields; multi-year-low leverage (~4.7x); a $373M deal locking long-lead electrical gear. 2026 core FFO guide $8.00–8.10; consensus $218 (+15%); 2.7% yield.
- **Watch:** ~24x forward FFO (top of band) for ~9% near-term growth, with a ~19-month booking-to-cash lag. Accumulate on rate-driven weakness; higher-for-longer is the key REIT headwind.

Peers

- **Equinix (EQIX) - Wealth Management Buy:** highest-quality digital-infra name; interconnection moat, 51% EBITDA margin, ~25x AFFO; consensus +12%, 1.8% yield. Extended near-term.
- **Iron Mountain (IRM) - Hold:** datacenter revenue +47% but ~13% of the company; a transition story with leverage. Consensus now sits ~flat to spot after a sharp run.

The Nordic gap

No clean public pure-play Nordic operator exists: atNorth, Bulk and EcoDataCenter are private; Green Mountain sits inside a diversified listed parent; Bahnhof is a listed ISP. Nordic thermal advantage has no clean listed pure-play and is only investable indirectly, at low purity, through the same REIT names that carry the Sun Belt exposure; that is a real limitation of the second leg and is why it is held as a thesis rather than sized as a position.

7 · Silicon & hyperscalers - hold the winners

Fundamentally still firing; positionally the most crowded trade on the Street. Keep quality, don't add beta, route the marginal dollar upstream into power.

On the apparent contradiction: several silicon names screen with large consensus upside (NVDA +34%, MSFT +25%). That is the Street's target, not my conviction - the “stop adding” call is about position size in the most-crowded trade, not a claim the names are overvalued. Hold the core; let the marginal dollar go to the scarcer power layer.

- **NVIDIA - Wealth Management Buy, Trading Hold:** Q1 FY27 revenue $81.6B (+85%), datacenter $75.2B (+92%), ~75% gross margin; next quarter guided ~$91B assuming ~zero China. At ~21x forward with consensus +34%, it is not the expensive name - but it is the epicenter of the crowded trade. Own it; don't add here.
- **TSMC (TSM) - Buy / Buy:** prints regardless of which chip wins; ~66% gross margin, FY26 growth guided >30%, ~22x. Best risk/reward; preferred silicon holding. ADR carries TWD exposure.
- **Alphabet (GOOGL) - Buy / Buy:** ~24x, cloud +63%, ~$462B backlog, TPU optionality; explicitly compute-constrained. Best risk/reward in mega-cap.
- **Amazon / Microsoft - Wealth Management Buy:** AWS reaccelerating (AMZN consensus +18%); Azure +40% but capacity-gated (MSFT consensus +25%). Own both; patient tactically.
- **Broadcom - Wealth Management Buy:** custom-ASIC hedge (~70% share); quality but only +7% to consensus into the June print.
- **Marvell - Trading Trim:** best fundamentals, most stretched (~51x, ~+1% to consensus). The semis short/hedge leg.

8 · Power producers - selective

The nuclear-IPP trade already ran - 2–4x in 2024–25, now de-rating (Constellation ~−22% off its high). Narrative intact; the re-rating is done. Stock-pick on contracts. Note the wide consensus upside here (VST +48%, CEG +41%) reflects targets set before the de-rate - treat as a range.

- **Talen (TLN) - Buy / Buy:** purest, most contract-levered nuclear-AI name; ~$18B / 1,920 MW Amazon PPA; the only one of the trio not de-rating. Consensus +18%. Single-asset concentration is the risk.
- **Vistra (VST) - Buy / Buy:** cheapest IPP multiple (~25x but lower on EV/EBITDA); 20-yr Meta (~2,600 MW) and AWS Comanche Peak (~1,200 MW) PPAs. High beta to power prices.
- **Constellation (CEG) - Wealth Management Buy:** largest US nuclear fleet; closed Calpine (~$16.4B equity) in Jan '26. De-rated; needs catalyst clearance (TMI restart, mandated gas divestiture).

GOES / electrical steel

A real choke point with no clean pure-play. The only US producer is Cleveland-Cliffs - a levered, loss-making turnaround, ownable only as a tariff trade. Cleaner to own the scarcity through the transformer OEMs that consume it (Section 5).

9 · The underweight / hedge leg

These fund the long book. Most over-owned, priced for perfection, or structurally mis-fit to the thesis.

Name	Why it's the worst risk/reward	Action
Vertiv (VRT)	~50x forward; cooling is the most cyclical, capex-sensitive link; orders +252% is a peak to fade	**Underweight**
Marvell (MRVL)	~51x and ~+1% to consensus; best fundamentals, most stretched - the semis hedge leg	**Trim**
Generic wholesale REITs	12–15x FFO without an interconnection moat is a value trap; rate-sensitive	**Avoid**
APLD / CoreWeave	Single-customer concentration + cash burn; leveraged calls, not core holdings	**Size tiny**
Cleveland-Cliffs (CLF)	Right bottleneck, wrong vehicle - distressed, levered, loss-making	**Avoid**

Note on GE Vernova: not on this list. It is a Buy on the franchise; the only caution is entry at ~69x - a timing call, addressed in Section 5.

10 · Top risks

Ranked by damage. Risks 1, 2 and 4 are correlated - an AI-ROI scare is also a capex-cut and a crowded-unwind, landing with no rate-cut cushion. Odds bands: High above 50%, Medium 25 to 50%, Low under 25%, over the 18-to-24-month horizon.

#	Risk	Odds	Hits hardest
1	**Hyperscaler capex cut / AI-ROI disappointment.** Capex growing ~3x faster than AI revenue.	Med-High	VRT, GEV; then REITs, IPPs
2	**Higher-for-longer macro.** Oil-shock CPI keeps the Fed frozen; 10-yr ~4.45%.	High	DLR/EQIX, then GEV/VRT
3	**Bottleneck moves behind-the-meter. On-site gas now, SMRs post-2030. But both still need step-up transformers and switchgear from the same OEMs, so the chokepoint migrates rather than clears.**	High	Grid-transformer names; helps IPPs
4	**Crowded-positioning unwind.** Six buyers fund the whole chain.	Medium	All names at once
5	**2027–28 capacity wave.** New OEM capacity normalizes pricing power.	Med (later)	Eaton, GEV-grid

What does not break it near-term: solid-state transformers (2030+), Chinese GOES (US tariff-blocked), FERC reform (multi-year). The thesis is good for ~18–24 months; the 2027–28 capacity wave is the natural exit signal.

11 · Illustrative positioning

Relative emphasis, not a portfolio allocation, and not personalized advice.

Sleeve	Emphasis	Expression
Power & grid	**Overweight**	Core ENR + GEV (add on pullbacks); ETN/PWR satellites
Silicon (quality)	**Market-weight**	TSM + GOOGL core; hold NVDA; don't add beta
Datacenter REITs	**Market-weight**	DLR + EQIX, accumulate on rate-driven weakness
Power producers	**Selective**	TLN/VST on contracts; not sector beta
Cooling / high-beta	**Underweight**	Short/trim VRT, MRVL; avoid generic wholesale

The discipline: almost every name here is the same crowded, capex-derivative bet under a different ticker. The thesis can be right on the shortage and still draw down 30–40% in a demand-side or macro air pocket. Size for that correlation, keep dry powder, favor the scarce layer (power) over the crowded one (silicon).

Sources and as-of. Prices, multiples, and dividend yields in Section II are as of the June 1, 2026 close (Hitachi as of May 8) and will drift; treat them as a dated snapshot. Consensus price targets are third-party aggregator medians, not house estimates. Company-specific figures (segment backlogs, bookings, lease and PPA terms, hyperscaler capex) are drawn from the most recent quarterly releases, earnings-call transcripts, and investor presentations of each named company;

macro positioning figures (Fed funds, CPI, manager-survey crowding) are from the corresponding official releases and bank fund-manager surveys as of the same date. Figures should be re-verified before the note is acted on or recirculated.

Disclaimer. For discussion and educational purposes only; not investment advice, not a research report under any securities regulation, and not an offer or solicitation. The author is an independent researcher, not a registered investment adviser or broker-dealer; nothing here is tailored to any individual's circumstances. Ratings and price-target reads are opinions, carry forward-looking uncertainty, and may be wrong; consensus targets are third-party aggregates, not the author's. Prices and multiples are from third-party sources believed reliable but not guaranteed, as of June 1, 2026 (Hitachi May 8); several names carry wide target dispersion. Foreign/ADR lines carry FX risk. Markets move - verify before acting. Consult a licensed adviser and do your own diligence. No liability is accepted for reliance on this note. Prepared by an independent researcher; figures compiled from public sources and not independently audited. The authorship and AI-assistance disclosure for this work is in Section I.